

Clang & LLVM: How they can improve your life as a developer

Tilmann Scheller

- Introduction
- LLVM Overview
- Clang
- Performance
- Summary



- Mature, production-quality compiler framework
- Modular architecture
- Heavily optimizing static and dynamic compiler
- Supports all major architectures (x86, ARM, MIPS, PowerPC, ...)
- Powerful link-time optimizations (LTO)
- Permissive license (BSD-like)

- Clang

C/C++/Objective C frontend and static analyzer

- LLDB

Next generation debugger leveraging the LLVM libraries, e.g. the Clang expression parser

- lld

Framework for creating linkers, will make Clang independent of the system linker in the future

- Polly

Polyhedral optimizer for LLVM, e.g. high-level loop optimizations and data-locality optimizations

Which companies are contributing?


SAMSUNG



Google




ARM[®]



QUALCOMM[®]

SONY



COMPUTER
ENTERTAINMENT[®]



NVIDIA[®]




CRAY

AMD



AZUL
SYSTEMS[®]



- WebKit FTL JIT
- Rust
- Android (NDK, ART, RenderScript)
- Portable NativeClient (PNaCl)
- Majority of OpenCL implementations based on Clang/LLVM
- CUDA, RenderScript
- LLVM on Linux: LLVMLinux, LLVMpipe (software rasterizer in Mesa), Radeon R300-R900 drivers in Mesa



OpenCL





FreeBSD®

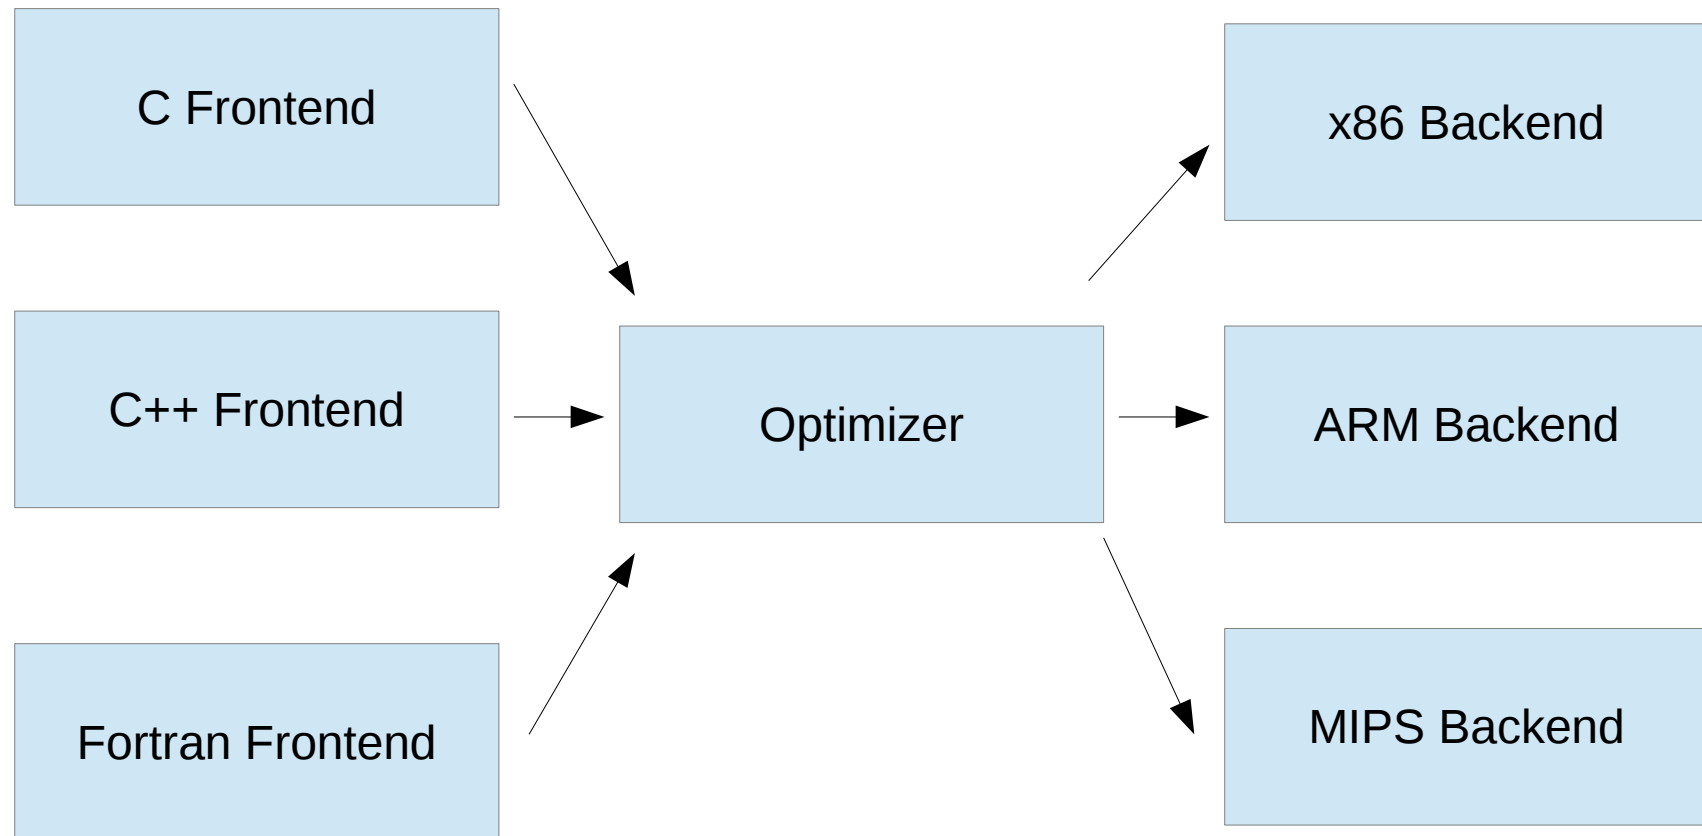
- Default compiler on OS X
- Default compiler on FreeBSD
- Default compiler for native applications on Tizen
- Default compiler on OpenMandriva Lx starting with the next release (2015.0)
- Debian experimenting with Clang as an additional compiler (94.3% of ~20k packages successfully build with Clang 3.5)
- Android NDK ships Clang



debian

TIZEN™ 

- LLVM IR
- Scalar optimizations
- Interprocedural optimizations
- Auto-vectorizer (BB, Loop and SLP)
- Profile-guided optimizations



- Many steps involved in the translation from C source code to machine code:
 - Frontend:
 - Lexing, Parsing, AST construction
 - Translation to LLVM IR
 - Middle-end
 - Target-independent optimizations (Analyses & Transformations)
 - Backend:
 - Translation into a DAG
 - Instruction selection: Pattern matching on the DAG
 - Instruction scheduling: Assigning an order of execution
 - Register allocation: Trying to reduce memory traffic

- The representation of the middle-end
- The majority of optimizations is done at LLVM IR level
- Low-level representation which carries type information
- RISC-like three-address code in static single assignment form with an infinite number of virtual registers
- Three different formats: bitcode (compact on-disk format), in-memory representation and textual representation (LLVM assembly language)

- Arithmetic: add, sub, mul, udiv, sdiv, ...

`%tmp = add i32 %indvar, -512`

- Logical operations: shl, lshr, ashr, and, or, xor

`%shr21 = ashr i32 %mul20, 8`

- Memory access: load, store, alloca, getelementptr

`%tmp3 = load i64* %tmp2`

- Comparison: icmp, select

`%cmp12 = icmp slt i32 %add, 1024`

- Control flow: call, ret, br, switch, ...


`call void @foo(i32 %phitmp)`

- Types: integer, floating point, vector, structure, array, ...

`i32, i342, double, <4 x float>, {i8, <2 x i16>}, [40 x i32]`

- Part of the backend
- Domain specific language to describe the instruction set, register file, calling conventions (TableGen)
- Pattern matcher is generated automatically
- Backend is a mix of C++ and TableGen
- Usually generates assembly code, direct machine code emission is also possible

```
zx = zy = zx2 = zy2 = 0;  
for (; iter < max_iter && zx2 + zy2 < 4; iter++) {  
    zy = 2 * zx * zy + y;  
    zx = zx2 - zy2 + x;  
    zx2 = zx * zx;  
    zy2 = zy * zy;  
}
```



```
zx = zy = zx2 = zy2 = 0;
for (; iter < max_iter && zx2 + zy2 < 4; iter++) {
    zy = 2 * zx * zy + y;
    zx = zx2 - zy2 + x;
    zx2 = zx * zx;
    zy2 = zy * zy;
}
```

loop:

```
%zy2.06 = phi double [ %8, %loop ], [ 0.000000e+00, %preheader ]
%zx2.05 = phi double [ %7, %loop ], [ 0.000000e+00, %preheader ]
%zy.04 = phi double [ %4, %loop ], [ 0.000000e+00, %preheader ]
%zx.03 = phi double [ %6, %loop ], [ 0.000000e+00, %preheader ]
%iter.02 = phi i32 [ %9, %loop ], [ 0, %.lr.ph.preheader ]
%2 = fmul double %zx.03, 2.000000e+00
%3 = fmul double %2, %zy.04
%4 = fadd double %3, %y
%5 = fsub double %zx2.05, %zy2.06
%6 = fadd double %5, %x
%7 = fmul double %6, %6
%8 = fmul double %4, %4
%9 = add i32 %iter.02, 1
%10 = icmp ult i32 %9, %max_iter
%11 = fadd double %7, %8
%12 = fcmp olt double %11, 4.000000e+00
%or.cond = and i1 %10, %12
br i1 %or.cond, label %loop, label %loopexit
```

loop:

```
// zx = zy = zx2 = zy2 = 0;
%zy2.06 = phi double [ %8, %loop ], [ 0.000000e+00, %preheader ]
%zx2.05 = phi double [ %7, %loop ], [ 0.000000e+00, %preheader ]
%zy.04 = phi double [ %4, %loop ], [ 0.000000e+00, %preheader ]
%zx.03 = phi double [ %6, %loop ], [ 0.000000e+00, %preheader ]
%iter.02 = phi i32 [ %9, %loop ], [ 0, %preheader ]
// zy = 2 * zx * zy + y;
%2 = fmul double %zx.03, 2.000000e+00
%3 = fmul double %2, %zy.04
%4 = fadd double %3, %y
// zx = zx2 - zy2 + x;
%5 = fsub double %zx2.05, %zy2.06
%6 = fadd double %5, %x
// zx2 = zx * zx;
%7 = fmul double %6, %6
// zy2 = zy * zy;
%8 = fmul double %4, %4
// iter++
%9 = add i32 %iter.02, 1
// iter < max_iter
%10 = icmp ult i32 %9, %max_iter
// zx2 + zy2 < 4
%11 = fadd double %7, %8
%12 = fcmp olt double %11, 4.000000e+00
// &&
%or.cond = and i1 %10, %12
br i1 %or.cond, label %loop, label %loopexit
```

```
zx = zy = zx2 = zy2 = 0;
for (;
    iter < max_iter
    && zx2 + zy2 < 4;
    iter++) {
    zy = 2 * zx * zy + y;
    zx = zx2 - zy2 + x;
    zx2 = zx * zx;
    zy2 = zy * zy;
}
```



```

zx = zy = zx2 = zy2 = 0;
for (;
    iter < max_iter
    && zx2 + zy2 < 4;
    iter++) {
    zy = 2 * zx * zy + y;
    zx = zx2 - zy2 + x;
    zx2 = zx * zx;
    zy2 = zy * zy;
}

```

.LBB0_2:

```

@ d17 = 2 * zx
vadd.f64          d17, d12, d12
@ iter < max_iter
cmp              r1, r0
@ d17 = (2 * zx) * zy
vmul.f64          d17, d17, d11
@ d18 = zx2 - zy2
vsub.f64          d18, d10, d8
@ d12 = (zx2 - zy2) + x
vadd.f64          d12, d18, d0
@ d11 = (2 * zx * zy) + y
vadd.f64          d11, d17, d9
@ zx2 = zx * zx
vmul.f64          d10, d12, d12
@ zy2 = zy * zy
vmul.f64          d8, d11, d11
bhs              .LBB0_5

```

@ BB#3:

```

@ zx2 + zy2
vadd.f64          d17, d10, d8
@ iter++
adds             r1, #1
@ zx2 + zy2 < 4
vcmpe.f64         d17, d16
vmrs              APSR_nzcv, fpscr
bmi              .LBB0_2
b                .LBB0_5

```

- Goals:
 - Fast compile time
 - Low memory usage
 - GCC compatibility
 - Expressive diagnostics
- Several tools built on top of Clang:
 - Clang static analyzer
 - clang-format, clang-modernize, clang-tidy

```
[t@ws-520 examples]$ cat t1.c
int a[2][2] = { 0, 1 , 2, 3 };
```

```
[t@ws-520 examples]$ clang-3.5 -c -Wall t1.c
t1.c:1:17: warning: suggest braces around initialization of subobject [-Wmissing-braces]
int a[2][2] = { 0, 1 , 2, 3 };
               ^~~~
               {   }
t1.c:1:24: warning: suggest braces around initialization of subobject [-Wmissing-braces]
int a[2][2] = { 0, 1 , 2, 3 };
               ^~~~
               {   }

2 warnings generated.
```

```
[t@ws-520 examples]$ gcc-4.9 -c -Wall t1.c
t1.c:1:1: warning: missing braces around initializer [-Wmissing-braces]
  int a[2][2] = { 0, 1 , 2, 3 };
  ^
t1.c:1:1: warning: (near initialization for 'a[0]') [-Wmissing-braces]
```

```
[t@ws-520 examples]$ cat t2.cpp
class A {
    A(int _a, int _b) : a(_a, b(_b) {}

    int a, b;
}
```

```
[t@ws-520 examples]$ clang++-3.5 -c -Wall t2.cpp
```

```
t2.cpp:4:13: error: expected ')'
```

```
    int a, b;
           ^
```

```
t2.cpp:2:26: note: to match this '('
```

```
    A(int _a, int _b) : a(_a, b(_b) {}
                        ^
```

```
t2.cpp:5:2: error: expected ';' after class
```

```
    }
    ^
```

```
    ;
```

```
2 errors generated.
```

```
[t@ws-520 examples]$ g++-4.9 -c -Wall t2.cpp
```

```
t2.cpp:5:1: error: expected ';' after class definition
```

```
    }
    ^
```

```
t2.cpp: In constructor 'A::A(int, int)':
```

```
t2.cpp:2:25: error: class 'A' does not have any field named 'a'
```

```
    A(int _a, int _b) : a(_a, b(_b) {}
                        ^
```

```
t2.cpp:2:35: error: 'b' was not declared in this scope
```

```
    A(int _a, int _b) : a(_a, b(_b) {}
                        ^
```

```
t2.cpp:4:12: error: expected ')' at end of input
```

```
    int a, b;
           ^
```

```
t2.cpp:4:12: error: expected '{' at end of input
```

```
[t@ws-520 examples]$ cat t3.cpp
extern bool f(int n);

void g(int a, int b)
{
    if (f(a) == 2)
        f(b);
}
```

```
[t@ws-520 examples]$ clang++-3.5 -c -Wall t3.cpp
t3.cpp:5:12: warning: comparison of constant 2 with expression of type 'bool'
is always false [-Wtautological-constant-out-of-range-compare]
    if (f(a) == 2)
        ~~~~ ^ ~
1 warning generated.
```

```
[t@ws-520 examples]$ g++-4.9 -c -Wall t3.cpp
[t@ws-520 examples]$
```

```
[t@ws-520 examples]$ cat t4.c
void foo(char *str) {
    strcpy(str, "foo");
}
```

```
[t@ws-520 examples]$ clang-3.5 -c -Wall t4.c
t4.c:2:3: warning: implicitly declaring library function 'strcpy' with
      type 'char *(char *, const char *)'
    strcpy(str, "foo");
    ^
t4.c:2:3: note: include the header <string.h> or explicitly provide a
      declaration for 'strcpy'
1 warning generated.
```

```
[t@ws-520 examples]$ gcc-4.9 -c -Wall t4.c
t4.c: In function 'foo':
t4.c:2:3: warning: implicit declaration of function 'strcpy' [-Wimplicit-function-declaration]
    strcpy(str, "foo");
    ^
t4.c:2:3: warning: incompatible implicit declaration of built-in function 'strcpy'
```

```
[t@ws-520 examples]$ cat t5.c
#include <stdio.h>
void foo(void) {
    printf("%s %d", "Hello, world");
}
```

```
[t@ws-520 examples]$ clang-3.5 -c -Wall t5.c
t5.c:3:15: warning: more '%' conversions than data arguments [-Wformat]
    printf("%s %d", "Hello, world");
           ~^
1 warning generated.
```

```
[t@ws-520 examples]$ gcc-4.9 -c -Wall t5.c
t5.c: In function 'foo':
t5.c:3:3: warning: format '%d' expects a matching 'int' argument [-Wformat=]
    printf("%s %d", "Hello, world");
    ^
```

- Part of Clang
- Tries to find bugs without executing the program
- Slower than compilation
- False positives
- Source annotations
- Works best on C code
- Runs from the commandline (scan-build), web interface for results

- Core Checkers
- C++ Checkers
- Dead Code Checkers
- Security Checkers
- Unix Checkers

llvm-toolchain-snapshot-3.6~svn219049 - scan-build results - Google Chrome

llvm-toolchain-snapsh x

← → ↻ llvm.org/reports/scan-build/

llvm-toolchain-snapshot-3.6~svn219049 - scan-build results

User:	pbuilder@irill3
Working Directory:	/tmp/buildd/llvm-toolchain-snapshot-3.6~svn219049
Command Line:	/usr/bin/make -j 7 -C build-llvm VERBOSE=1 CLANG_VENDOR=Debian 'CXXFLAGS= -std=c++0x' 'LDFLAGS= -fuse-ld=gold' REQUIRES_RTTI=1 DEBUGMAKE=1
Clang Version:	Debian clang version 3.6.0-svn219049-1~exp1 (trunk) (based on LLVM 3.6.0)
Date:	Sat Oct 4 12:22:52 2014

Bug Summary

Bug Type	Quantity	Display?
All Bugs	228	<input checked="" type="checkbox"/>
API		
Argument with 'nonnull' attribute passed null	2	<input checked="" type="checkbox"/>
Dead store		
Dead assignment	53	<input checked="" type="checkbox"/>
Dead increment	7	<input checked="" type="checkbox"/>
Dead initialization	3	<input checked="" type="checkbox"/>
Logic error		
Array subscript is undefined	2	<input checked="" type="checkbox"/>
Assigned value is garbage or undefined	3	<input checked="" type="checkbox"/>
Branch condition evaluates to a garbage value	1	<input checked="" type="checkbox"/>
Called C++ object pointer is null	76	<input checked="" type="checkbox"/>
Called C++ object pointer is uninitialized	1	<input checked="" type="checkbox"/>
Called function pointer is null (null dereference)	3	<input checked="" type="checkbox"/>
Dereference of null pointer	23	<input checked="" type="checkbox"/>
Division by zero	6	<input checked="" type="checkbox"/>

Reports

Bug Group	Bug Type ▼	File	Function/Method	Line	Path Length	
Logic error	Assigned value is garbage or undefined	lib/Object/IObjectFile.cpp	moveSymbolNext	186	3	View Report
Logic error	Assigned value is garbage or undefined	lib/Support/ScaledNumber.cpp	toStringAPFloat	172	16	View Report
Logic error	Assigned value is garbage or undefined	lib/Target/X86/X86ISelLowering.cpp	operator()	8532	7	View Report
Dead store	Dead increment	lldb/tools/lldb-platform/lldb-platform.cpp	main	238	1	View Report
Dead store	Dead increment	lldb/source/Target/Process.cpp	WriteMemory	2536	1	View Report
Dead store	Dead increment	lldb/tools/lldb-mi/MICmnLLDBDebugSessionInfo.cpp	MIResponseFormBrkPtInfo	1253	1	View Report
Dead store	Dead increment	lldb/tools/lldb-mi/MICmnLLDBDebuggerHandleEvents.cpp	MIStoppedAtBreakPoint	1069	1	View Report
Dead store	Dead increment	lldb/tools/lldb-platform/lldb-platform.cpp	main	237	1	View Report
Dead store	Dead increment	lib/Target/ARM/Disassembler/ARMDisassembler.cpp	DecodeVLD4DupInstruction	2982	1	View Report
Dead store	Dead increment	lldb/source/Plugins/Process/POSIX/RegisterContextPOSIXProcessMonitor_mips64.cpp	WriteAllRegisterValues	217	1	View Report
Memory Error	Memory leak	build-llvm/tools/lldb/source/Interpreter/LLDBWrapPython.cpp	_wrap_SBTTarget_Launch__SWIG_0	40975	12	View Report
Memory Error	Memory leak	build-llvm/tools/lldb/source/Interpreter/LLDBWrapPython.cpp	_wrap_SBTTarget_LaunchSimple	41116	9	View Report
Memory Error	Memory leak	build-llvm/tools/lldb/source/Interpreter/LLDBWrapPython.cpp	_wrap_SBProcess_ReadCStringFromMemory	33539	9	View Report
Memory Error	Memory leak	clang/tools/c-index-test/c-index-test.c	perform_test_reparse_source	1598	28	View Report
Memory Error	Memory leak	clang/tools/c-index-test/c-index-test.c	find_file_includes_in	2448	7	View Report
Memory Error	Memory leak	build-llvm/tools/lldb/source/Interpreter/LLDBWrapPython.cpp	_wrap_SBProcess_ReadMemory	33405	9	View Report
Memory Error	Memory leak	unittests/IR/WaymarkTest.cpp	TestBody	47	2	View Report
Memory Error	Memory leak	build-llvm/tools/lldb/source/Interpreter/LLDBWrapPython.cpp	_wrap_SBProcess_RemoteLaunch	32437	9	View Report
Logic error	Result of operation is garbage or undefined	lib/CodeGen/AsmPrinter/AsmPrinter.cpp	EmitAlignment	1534	5	View Report
Logic error	Result of operation is garbage or undefined	lib/CodeGen/InlineSpiller.cpp	isSnippet	249	7	View Report
Logic error	Result of operation is garbage or undefined	tools/llvm-objdump/MachODump.cpp	SegInfo	2746	8	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/Mips/MipsAnalyzeImmediate.cpp	GetInstSeqLsSLL	44	25	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/AArch64/AArch64ISelDAGToDAG.cpp	SelectAddrModeUnscaled	648	5	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/NVPTX/NVPTXISelDAGToDAG.cpp	SelectBFE	4842	10	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/ARM/ARMISelLowering.cpp	PerformBFICombine	8576	4	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/ARM/ARMConstantIslandPass.cpp	UnknownPadding	69	8	View Report
Logic error	Result of operation is garbage or undefined	lib/Target/X86/X86ISelLowering.cpp	operator()	8476	8	View Report
Logic error	Result of operation is garbage or undefined	lib/CodeGen/ExecutionDepsFix.cpp	hasDomain	78	12	View Report
Logic error	Result of operation is garbage or undefined	lib/MC/MCAsmStreamer.cpp	EmitValueImpl	685	14	View Report

...

```
const SCEV *MaxBECOUNT = getCouldNotCompute();
```

Value stored to 'MaxBECOUNT' during its initialization is never read

```
if (isa<SCEVConstant>(BECOUNT))
```

```
    MaxBECOUNT = BECOUNT;
```

```
else
```

```
    MaxBECOUNT = computeBECOUNT(getConstant(MaxStart - MinEnd),  
                                getConstant(MinStride), false);
```

```
if (isa<SCEVCouldNotCompute>(MaxBECOUNT))
```

```
    MaxBECOUNT = BECOUNT;
```

```
return ExitLimit(BECOUNT, MaxBECOUNT, /*MustExit=*/true);
```

```
}
```

```
495 NestedNameSpecifierLocBuilder &
496 NestedNameSpecifierLocBuilder::
497 operator=(const NestedNameSpecifierLocBuilder &Other) {
498     Representation = Other.Representation;
499
500     if (Buffer && Other.Buffer && BufferCapacity >= Other.BufferSize) {
```

1 Taking false branch →

```
501     // Re-use our storage.
502     BufferSize = Other.BufferSize;
503     memcpy(Buffer, Other.Buffer, BufferSize);
504     return *this;
505 }
```

```
506
507 // Free our storage, if we have any.
508 if (BufferCapacity) {
```

2 → Taking true branch →

```
509     free(Buffer);
510
511     BufferCapacity = 0;
512 }
513 if (!Other.Buffer) {
```

4 → Taking false branch →

```
514     // Empty.
515     Buffer = nullptr;
516     BufferSize = 0;
517     return *this;
518 }
```

```
519
520 if (Other.BufferCapacity == 0) {
```

5 → Taking false branch →

```
521     // Shallow copy is okay.
522     Buffer = Other.Buffer;
523     BufferSize = Other.BufferSize;
524     return *this;
525 }
```

```
526
527 // Deep copy.
528 Append(Other.Buffer, Other.Buffer + Other.BufferSize, Buffer, BufferSize,
```

6 → Calling 'Append' →

```
529         BufferCapacity);
530     return *this;
531 }
```

```
435 namespace {
436     void Append(char *Start, char *End, char *&Buffer, unsigned &BufferSize,
437                unsigned &BufferCapacity) {
438         if (BufferSize + (End - Start) > BufferCapacity) {
439             // Reallocate the buffer.
440             unsigned NewCapacity
441                 = std::max((unsigned)(BufferSize + (End - Start)),
442                           : sizeof(void*) * 2,
443                           (unsigned)(BufferSize + (End - Start)));
444             char *NewBuffer = static_cast<char *>(malloc(NewCapacity));
445             memcpy(NewBuffer, Buffer, BufferSize);
446
447             if (BufferCapacity)
448                 free(Buffer);
449             Buffer = NewBuffer;
450             BufferCapacity = NewCapacity;
451         }
452
453         memcpy(Buffer + BufferSize, Start, End - Start);
454         BufferSize += End - Start;
455     }
```

7 → Taking true branch →

8 → '?' condition is false →

9 → Use of memory after it is freed

- Automatic formatting
- Developers waste time on formatting
- Supports different style guides
- Consistent coding style is important

- Detect bug prone coding patterns
- Enforce coding conventions
- Advocate modern and maintainable code
- Checks can be more expensive than compilation

- LLVM/Clang-based Sanitizer projects:
 - AddressSanitizer – Fast memory error detector
 - ThreadSanitizer – Detects data races
 - LeakSanitizer – Memory leak detector
 - MemorySanitizer – Detects reads of uninitialized variables
 - UBSanitizer – Detects undefined behavior

- LLVM 3.5 (released in September)
- Self-hosting on SPARC64 (Linux, FreeBSD)
- Integrated assembler and EHABI enabled by default on ARM
- Merging of the AArch64 backends completed
- Optimization reports
- Experimental support for C++17 features

- Better Windows support
- Improved debugging
- LTO parallelization
- Vectorization
- Profile-guided optimizations
 - Profiling infrastructure (sampling/instrumentation-based)
 - Analyses and transformations which take advantage of the gathered data

- Clang 3.5 and GCC 4.9.1 binaries both compiled with GCC 4.9.1 at -O2
- Clang 3.5 release build (-O3)
 - Clang 3.5: 6m46s
 - GCC 4.9.1: 9m56s
- Clang 3.5 debug build (-O0 -gsplit-dwarf)
 - Clang 3.5: 7m13s
 - GCC 4.9.1: 10m34s
- Clang ~46% faster in both builds!
- Measured on Fedora 20 (3.5GHz i7-4770K, 16GB RAM, 2TB HDD)

- Great compiler infrastructure
- Fast C/C++ compiler with expressive diagnostics
- Bug detection at compile time
- Automated formatting of code
- Detect memory bugs early with Sanitizers

- Visit llvm.org
- Distributions with Clang/LLVM packages:
 - Fedora
 - Debian/Ubuntu
 - openSUSE
 - ...

Thank you.

A decorative graphic at the bottom of the slide consisting of several overlapping, translucent blue waves that flow from left to right.