

## Attributes in Clang

### Introduction

### Function Attributes

**#pragma omp declare simd**  
**#pragma omp declare target**  
**\_Noreturn**  
**abi\_tag (gnu::abi\_tag)**  
**acquire\_capability (acquire\_shared\_capability, clang::acquire\_capability, clang::acquire\_shared\_capability)**  
**alloc\_align (gnu::alloc\_align)**  
**alloc\_size (gnu::alloc\_size)**  
**assert\_capability (assert\_shared\_capability, clang::assert\_capability, clang::assert\_shared\_capability)**  
**assume\_aligned (gnu::assume\_aligned)**  
**availability**  
**carries\_dependency**  
**convergent (clang::convergent)**  
**deprecated (gnu::deprecated)**  
**diagnose\_if**  
**disable\_tail\_calls (clang::disable\_tail\_calls)**  
**enable\_if**  
**external\_source\_symbol (clang::external\_source\_symbol)**  
**flatten (gnu::flatten)**  
**force\_align\_arg\_pointer (gnu::force\_align\_arg\_pointer)**  
**format (gnu::format)**  
**ifunc (gnu::ifunc)**  
**internal\_linkage (clang::internal\_linkage)**  
**interrupt (ARM)**  
**interrupt (AVR)**  
**interrupt (MIPS)**  
**kernel**  
**long\_call (gnu::long\_call, gnu::far)**  
**micromips (gnu::micromips)**

**no\_caller\_saved\_registers** (gnu::no\_caller\_saved\_registers)  
**no\_sanitize** (clang::no\_sanitize)  
**no\_sanitize\_address** (no\_address\_safety\_analysis, gnu::no\_address\_safety\_analysis, gnu::no\_sanitize\_address)  
**no\_sanitize\_memory**  
**no\_sanitize\_thread**  
**no\_split\_stack** (gnu::no\_split\_stack)  
**noalias**  
**nodiscard**, **warn\_unused\_result**, clang::warn\_unused\_result, gnu::warn\_unused\_result  
**noduplicate** (clang::noduplicate)  
**nomicrotips** (gnu::nomicrotips)  
**noreturn**  
**not\_tail\_called** (clang::not\_tail\_called)  
**nothrow** (gnu::nothrow)  
**objc\_boxable**  
**objc\_method\_family**  
**objc\_requires\_super**  
**objc\_runtime\_name**  
**objc\_runtime\_visible**  
**optnone** (clang::optnone)  
**overloadable**  
**release\_capability** (release\_shared\_capability, clang::release\_capability, clang::release\_shared\_capability)  
**short\_call** (gnu::short\_call, gnu::near)  
**signal** (gnu::signal)  
**target** (gnu::target)  
**try\_acquire\_capability** (try\_acquire\_shared\_capability, clang::try\_acquire\_capability, clang::try\_acquire\_shared\_capability)  
**xray\_always\_instrument** (clang::xray\_always\_instrument), **xray\_never\_instrument** (clang::xray\_never\_instrument),  
**xray\_log\_args** (clang::xray\_log\_args)  
**xray\_always\_instrument** (clang::xray\_always\_instrument), **xray\_never\_instrument** (clang::xray\_never\_instrument),  
**xray\_log\_args** (clang::xray\_log\_args)

#### Variable Attributes

**dlexport** (gnu::dlexport)  
**dllimport** (gnu::dllimport)  
**init\_seg**

**maybe\_unused, unused, gnu::unused**  
**nodebug** (gnu::nodebug)  
**noescape** (clang::noescape)  
**nosvm**  
**pass\_object\_size**  
**require\_constant\_initialization** (clang::require\_constant\_initialization)  
**section** (gnu::section, \_\_declspec(allocate))  
**swift\_context**  
**swift\_error\_result**  
**swift\_indirect\_result**  
**swiftcall**  
**thread**  
**tls\_model** (gnu::tls\_model)

#### Type Attributes

**\_\_single\_inheritance, \_\_multiple\_inheritance, \_\_virtual\_inheritance**  
**align\_value**  
**empty\_bases**  
**enum\_extensibility** (clang::enum\_extensibility)  
**flag\_enum**  
**layout\_version**  
**lto\_visibility\_public** (clang::lto\_visibility\_public)  
**novtable**  
**objc\_subclassing\_restricted**  
**selectany** (gnu::selectany)  
**transparent\_union** (gnu::transparent\_union)

#### Statement Attributes

**#pragma clang loop**  
**#pragma unroll, #pragma nounroll**  
**\_\_attribute\_\_((intel\_reqd\_sub\_group\_size))**  
**\_\_attribute\_\_((opencl\_unroll\_hint))**  
**\_\_read\_only, \_\_write\_only, \_\_read\_write** (read\_only, write\_only, read\_write)  
**fallthrough, clang::fallthrough**  
**suppress** (gsl::suppress)

#### Calling Conventions

**fastcall** (gnu::fastcall, \_\_fastcall, \_fastcall)  
**ms\_abi** (gnu::ms\_abi)  
**pcs** (gnu::pcs)  
**preserve\_all**  
**preserve\_most**  
**regcall** (gnu::regcall, \_\_regcall)  
**regparm** (gnu::regparm)  
**stdcall** (gnu::stdcall, \_\_stdcall, \_stdcall)  
**thiscall** (gnu::thiscall, \_\_thiscall, \_thiscall)  
**vectorcall** (\_\_vectorcall, \_vectorcall)

#### **AMD GPU Attributes**

**amdgpu\_flat\_work\_group\_size**  
**amdgpu\_num\_sgpr**  
**amdgpu\_num\_vgpr**  
**amdgpu\_waves\_per\_eu**

#### **Consumed Annotation Checking**

**callable\_when**  
**consumable**  
**param\_typestate**  
**return\_typestate**  
**set\_typestate**  
**test\_typestate**

#### **Type Safety Checking**

**argument\_with\_type\_tag**  
**pointer\_with\_type\_tag**  
**type\_tag\_for\_datatype**

#### **OpenCL Address Spaces**

**constant** (\_\_constant)  
**generic** (\_\_generic)  
**global** (\_\_global)  
**local** (\_\_local)  
**private** (\_\_private)

#### **Nullability Attributes**

**\_Nonnull**

**\_Null\_unspecified**  
**\_Nullable**  
**nonnull (gnu::nonnull)**  
**returns\_nonnull (gnu::returns\_nonnull)**

## Introduction

This page lists the attributes currently supported by Clang.

## Function Attributes

### #pragma omp declare simd

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
					X	

The *declare simd* construct can be applied to a function to enable the creation of one or more versions that can process multiple arguments using SIMD instructions from a single invocation in a SIMD loop. The *declare simd* directive is a declarative directive. There may be multiple *declare simd* directives for a function. The use of a *declare simd* construct on a function enables the creation of SIMD versions of the associated function that can be used to process multiple arguments from a single invocation from a SIMD loop concurrently. The syntax of the *declare simd* construct is as follows:

```
#pragma omp declare simd [clause[, clause] ...] new-line
#pragma omp declare simd [clause[, clause] ...] new-line [...] function
definition or declaration
```

#pragma omp declare simd [clause[, clause] ...] new-line [#pragma omp declare simd [clause[, clause] ...] new-line] [...] function definition or declaration

where clause is one of the following:

simdlen(length) linear(argument-list[:constant-linear-step]) aligned(argument-list[:alignment]) uniform(argument-list) inbranch  
notinbranch

## #pragma omp declare target

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
					X	

The *declare target* directive specifies that variables and functions are mapped to a device for OpenMP offload mechanism.

The syntax of the declare target directive is as follows:

```
#pragma omp declare target new-line declarations-definition-seq #pragma omp end declare target new-line
```

#pragma omp declare target new-line declarations-definition-seq #pragma omp end declare target new-line

## \_Noreturn

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

A function declared as `_Noreturn` shall not return to its caller. The compiler will generate a diagnostic for a function declared as `_Noreturn` that appears to be capable of returning to its caller.

## abi\_tag (gnu::abi\_tag)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `abi_tag` attribute can be applied to a function, variable, class or inline namespace declaration to modify the mangled name of the entity. It gives the ability to distinguish between different versions of the same entity but with different ABI versions supported. For example, a newer version of a class could have a different set of data members and thus have a different size. Using the `abi_tag` attribute, it is possible to have different mangled names for a global variable of the class type. Therefor, the old code could keep using the old manged name and the new code will use the new mangled name with tags.

### **acquire\_capability (acquire\_shared\_capability, clang::acquire\_capability, clang::acquire\_shared\_capability)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Marks a function as acquiring a capability.

### **alloc\_align (gnu::alloc\_align)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Use `__attribute__((alloc_align(<alignment>)))` on a function declaration to specify that the return value of the function (which must be a pointer type) is at least as aligned as the value of the indicated parameter. The parameter is given by its index in the list of formal parameters; the first parameter has index 1 unless the function is a C++ non-static member function, in which case the first parameter has index 2 to account for the implicit this parameter.

```
// The returned pointer has the alignment specified by the first parameter.
void *a(size_t align) __attribute__((alloc_align(1)));
```

```
// The returned pointer has the alignment specified by the second parameter.
```

```
void *b(void *v, size_t align) __attribute__((alloc_align(2)));
```

```
// The returned pointer has the alignment specified by the second visible
```

```
// parameter, however it must be adjusted for the implicit 'this' parameter.
```

```
void *Foo::b(void *v, size_t align) __attribute__((alloc_align(3)));
```

Note that this attribute merely informs the compiler that a function always returns a sufficiently aligned pointer. It does not cause the compiler to emit code to enforce that alignment. The behavior is undefined if the returned pointer is not sufficiently aligned.

## alloc\_size (gnu::alloc\_size)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `alloc_size` attribute can be placed on functions that return pointers in order to hint to the compiler how many bytes of memory will be available at the returned pointer. `alloc_size` takes one or two arguments.

`alloc_size(N)` implies that argument number N equals the number of available bytes at the returned pointer.

`alloc_size(N, M)` implies that the product of argument number N and argument number M equals the number of available bytes at the returned pointer.

Argument numbers are 1-based.

An example of how to use `alloc_size`

```
void *my_malloc(int a) __attribute__((alloc_size(1)));
void *my_calloc(int a, int b) __attribute__((alloc_size(1, 2)));

int main() {
    void *const p = my_malloc(100);
    assert(__builtin_object_size(p, 0) == 100);
    void *const a = my_calloc(20, 5);
}
```



```
assert(__builtin_object_size(a, 0) == 100);
}
```

### Note

This attribute works differently in clang than it does in GCC. Specifically, clang will only trace `const` pointers (as above); we give up on pointers that are not marked as `const`. In the vast majority of cases, this is unimportant, because LLVM has support for the `alloc_size` attribute. However, this may cause mildly unintuitive behavior when used with other attributes, such as `enable_if`.

## **assert\_capability (assert\_shared\_capability, clang::assert\_capability, clang::assert\_shared\_capability)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Marks a function that dynamically tests whether a capability is held, and halts the program if it is not held.

## **assume\_aligned (gnu::assume\_aligned)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Use `__attribute__((assume_aligned(<alignment>[,<offset>])))` on a function declaration to specify that the return value of the function (which must be a pointer type) has the specified offset, in bytes, from an address with the specified alignment. The offset is taken to be zero if omitted.

```
// The returned pointer value has 32-byte alignment.
```

```
void *a() __attribute__((assume_aligned (32)));
```

```
// The returned pointer value is 4 bytes greater than an address having
```

```
// 32-byte alignment.
void *b() __attribute__((assume_aligned (32, 4)));
```

Note that this attribute provides information to the compiler regarding a condition that the code already ensures is true. It does not cause the compiler to enforce the provided alignment assumption.

## availability

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `availability` attribute can be placed on declarations to describe the lifecycle of that declaration relative to operating system versions. Consider the function declaration for a hypothetical function `f`:

```
void f(void) __attribute__((availability(macos,introduced=10.4,deprecated=10.6,obsoleted=10.7)));
```

The `availability` attribute states that `f` was introduced in macOS 10.4, deprecated in macOS 10.6, and obsoleted in macOS 10.7. This information is used by Clang to determine when it is safe to use `f`: for example, if Clang is instructed to compile code for macOS 10.5, a call to `f()` succeeds. If Clang is instructed to compile code for macOS 10.6, the call succeeds but Clang emits a warning specifying that the function is deprecated. Finally, if Clang is instructed to compile code for macOS 10.7, the call fails because `f()` is no longer available.

The `availability` attribute is a comma-separated list starting with the platform name and then including clauses specifying important milestones in the declaration's lifetime (in any order) along with additional information. Those clauses can be:

`introduced=version`

The first version in which this declaration was introduced.

`deprecated=version`

The first version in which this declaration was deprecated, meaning that users should migrate away from this API.

`obsoleted=version`

The first version in which this declaration was obsoleted, meaning that it was removed completely and can no longer be used.

unavailable

This declaration is never available on this platform.

message=*string-literal*

Additional message text that Clang will provide when emitting a warning or error about use of a deprecated or obsoleted declaration. Useful to direct users to replacement APIs.

replacement=*string-literal*

Additional message text that Clang will use to provide Fix-It when emitting a warning about use of a deprecated declaration. The Fix-It will replace the deprecated declaration with the new declaration specified.

Multiple availability attributes can be placed on a declaration, which may correspond to different platforms. Only the availability attribute with the platform corresponding to the target platform will be used; any others will be ignored. If no availability attribute specifies availability for the current target platform, the availability attributes are ignored. Supported platforms are:

ios

Apple's iOS operating system. The minimum deployment target is specified by the `-mios-version-min=*version*` or `-miphoneos-version-min=*version*` command-line arguments.

macos

Apple's macOS operating system. The minimum deployment target is specified by the `-mmacosx-version-min=*version*` command-line argument. `macosx` is supported for backward-compatibility reasons, but it is deprecated.

tvos

Apple's tvOS operating system. The minimum deployment target is specified by the `-mtvos-version-min=*version*` command-line argument.

watchos

Apple's watchOS operating system. The minimum deployment target is specified by the `-mwatchos-version-min=*version*` command-line argument.

A declaration can typically be used even when deploying back to a platform version prior to when the declaration was introduced. When this happens, the declaration is **weakly linked**, as if the `weak_import` attribute were added to the declaration. A weakly-linked declaration may or may not be present at run-time, and a program can determine whether the declaration is present by checking whether the address of that declaration is non-NULL.

The flag `strict` disallows using API when deploying back to a platform version prior to when the declaration was introduced. An attempt to use such API before its introduction causes a hard error. Weakly-linking is almost always a better API choice, since it allows users to query

availability at runtime.

If there are multiple declarations of the same entity, the availability attributes must either match on a per-platform basis or later declarations must not have availability attributes for that platform. For example:

```
void g(void) __attribute__((availability(macos,introduced=10.4)));
void g(void) __attribute__((availability(macos,introduced=10.4))); // okay, matches
void g(void) __attribute__((availability(ios,introduced=4.0))); // okay, adds a new platform
void g(void); // okay, inherits both macos and ios availability from above.
void g(void) __attribute__((availability(macos,introduced=10.5))); // error: mismatch
```

When one method overrides another, the overriding method can be more widely available than the overridden method, e.g.,:

```
@interface A
- (id)method __attribute__((availability(macos,introduced=10.4)));
- (id)method2 __attribute__((availability(macos,introduced=10.4)));
@end

@interface B : A
- (id)method __attribute__((availability(macos,introduced=10.3))); // okay: method moved into base class later
- (id)method __attribute__((availability(macos,introduced=10.5))); // error: this method was available via the base class in 10.4
@end
```

Starting with the macOS 10.12 SDK, the `API_AVAILABLE` macro from `<os/availability.h>` can simplify the spelling:

```
@interface A
- (id)method API_AVAILABLE(macos(10.11));
- (id)otherMethod API_AVAILABLE(macos(10.11), ios(11.0));
@end
```

Also see the documentation for **@available**

**carries\_dependency**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `carries_dependency` attribute specifies dependency propagation into and out of functions.

When specified on a function or Objective-C method, the `carries_dependency` attribute means that the return value carries a dependency out of the function, so that the implementation need not constrain ordering upon return from that function. Implementations of the function and its caller may choose to preserve dependencies instead of emitting memory ordering instructions such as fences.

Note, this attribute does not change the meaning of the program, but may result in generation of more efficient code.

### convergent (clang::convergent)

#### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `convergent` attribute can be placed on a function declaration. It is translated into the LLVM `convergent` attribute, which indicates that the call instructions of a function with this attribute cannot be made control-dependent on any additional values.

In languages designed for SPMD/SIMT programming model, e.g. OpenCL or CUDA, the call instructions of a function with this attribute must be executed by all work items or threads in a work group or sub group.

This attribute is different from `noduplicate` because it allows duplicating function calls if it can be proved that the duplicated function calls are not made control-dependent on any additional values, e.g., unrolling a loop executed by all work items.

Sample usage: .. code-block:: c

```
void convfunc(void) __attribute__((convergent)); // Setting it as a C++11 attribute is also valid in a C++ program. // void
convfunc(void) [[clang::convergent]];
```

### deprecated (gnu::deprecated)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X	X	X			

The `deprecated` attribute can be applied to a function, a variable, or a type. This is useful when identifying functions, variables, or types that are expected to be removed in a future version of a program.

Consider the function declaration for a hypothetical function `f`:

```
void f(void) __attribute__((deprecated("message", "replacement")));
```

When spelled as `__attribute__((deprecated))`, the deprecated attribute can have two optional string arguments. The first one is the message to display when emitting the warning; the second one enables the compiler to provide a Fix-It to replace the deprecated name with a new name. Otherwise, when spelled as `[[gnu::deprecated]]` or `[[deprecated]]`, the attribute can have one optional string argument which is the message to display when emitting the warning.

## diagnose\_if

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

The `diagnose_if` attribute can be placed on function declarations to emit warnings or errors at compile-time if calls to the attributed function meet certain user-defined criteria. For example:

```
void abs(int a)
__attribute__((diagnose_if(a >= 0, "Redundant abs call", "warning")));
void must_abs(int a)
__attribute__((diagnose_if(a >= 0, "Redundant abs call", "error")));

int val = abs(1); // warning: Redundant abs call
int val2 = must_abs(1); // error: Redundant abs call
```

```
int val3 = abs(val);
int val4 = must_abs(val); // Because run-time checks are not emitted for
                        // diagnose_if attributes, this executes without
                        // issue.
```

`diagnose_if` is closely related to `enable_if`, with a few key differences:

Overload resolution is not aware of `diagnose_if` attributes: they're considered only after we select the best candidate from a given candidate set.

Function declarations that differ only in their `diagnose_if` attributes are considered to be redeclarations of the same function (not overloads). If the condition provided to `diagnose_if` cannot be evaluated, no diagnostic will be emitted.

Otherwise, `diagnose_if` is essentially the logical negation of `enable_if`.

As a result of bullet number two, `diagnose_if` attributes will stack on the same function. For example:

```
int foo() __attribute__((diagnose_if(1, "diag1", "warning")));
int foo() __attribute__((diagnose_if(1, "diag2", "warning")));

int bar = foo(); // warning: diag1
               // warning: diag2
int (*fooptr)(void) = foo; // warning: diag1
                        // warning: diag2

constexpr int supportsAPILevel(int N) { return N < 5; }
int baz(int a)
    __attribute__((diagnose_if(!supportsAPILevel(10),
                             "Upgrade to API level 10 to use baz", "error")));
int baz(int a)
    __attribute__((diagnose_if(!a, "0 is not recommended.", "warning")));

int (*bazptr)(int) = baz; // error: Upgrade to API level 10 to use baz
int v = baz(0); // error: Upgrade to API level 10 to use baz
```

Query for this feature with `__has_attribute(diagnose_if)`.

## disable\_tail\_calls (clang::disable\_tail\_calls)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `disable_tail_calls` attribute instructs the backend to not perform tail call optimization inside the marked function.

For example:

```
int callee(int);

int foo(int a) __attribute__((disable_tail_calls)) {
    return callee(a); // This call is not tail-call optimized.
}
```

Marking virtual functions as `disable_tail_calls` is legal.

```
int callee(int);

class Base {
public:
    [[clang::disable_tail_calls]] virtual int foo1() {
        return callee(); // This call is not tail-call optimized.
    }
};

class Derived1 : public Base {
public:
    int foo1() override {
        return callee(); // This call is tail-call optimized.
    }
};
```

## enable\_if



## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

**Note**

Some features of this attribute are experimental. The meaning of multiple `enable_if` attributes on a single declaration is subject to change in a future version of clang. Also, the ABI is not standardized and the name mangling may change in future versions. To avoid that, use asm labels.

The `enable_if` attribute can be placed on function declarations to control which overload is selected based on the values of the function's arguments. When combined with the `overloadable` attribute, this feature is also available in C.

```
int isdigit(int c);
int isdigit(int c) __attribute__((enable_if(c <= -1 || c > 255, "chosen when 'c' is out of range"))) __attribute__((unavailable("'c' must have the value of an unsigned char or EOF")));

void foo(char c) {
    isdigit(c);
    isdigit(10);
    isdigit(-10); // results in a compile-time error.
}
```

The `enable_if` attribute takes two arguments, the first is an expression written in terms of the function parameters, the second is a string explaining why this overload candidate could not be selected to be displayed in diagnostics. The expression is part of the function signature for the purposes of determining whether it is a redeclaration (following the rules used when determining whether a C++ template specialization is ODR-equivalent), but is not part of the type.

The `enable_if` expression is evaluated as if it were the body of a bool-returning constexpr function declared with the arguments of the function it is being applied to, then called with the parameters at the call site. If the result is false or could not be determined through constant expression evaluation, then this overload will not be chosen and the provided string may be used in a diagnostic if the compile fails as a result.

Because the `enable_if` expression is an unevaluated context, there are no global state changes, nor the ability to pass information from the `enable_if` expression to the function body. For example, suppose we want calls to `strlen(strbuf, maxlen)` to resolve to `strlen_chk(strbuf, maxlen, size of strbuf)` only if the size of `strbuf` can be determined:

```
__attribute__((always_inline))
static inline size_t strlen(const char *s, size_t maxlen)
__attribute__((overloadable))
__attribute__((enable_if(__builtin_object_size(s, 0) != -1))),
    "chosen when the buffer size is known but 'maxlen' is not")))
{
    return strlen_chk(s, maxlen, __builtin_object_size(s, 0));
}
```

Multiple `enable_if` attributes may be applied to a single declaration. In this case, the `enable_if` expressions are evaluated from left to right in the following manner. First, the candidates whose `enable_if` expressions evaluate to false or cannot be evaluated are discarded. If the remaining candidates do not share ODR-equivalent `enable_if` expressions, the overload resolution is ambiguous. Otherwise, `enable_if` overload resolution continues with the next `enable_if` attribute on the candidates that have not been discarded and have remaining `enable_if` attributes. In this way, we pick the most specific overload out of a number of viable overloads using `enable_if`.

```
void f() __attribute__((enable_if(true, ""))); // #1
void f() __attribute__((enable_if(true, ""))) __attribute__((enable_if(true, ""))); // #2

void g(int i, int j) __attribute__((enable_if(i, ""))); // #1
void g(int i, int j) __attribute__((enable_if(j, ""))) __attribute__((enable_if(true, ""))); // #2
```

In this example, a call to `f()` is always resolved to `#2`, as the first `enable_if` expression is ODR-equivalent for both declarations, but `#1` does not have another `enable_if` expression to continue evaluating, so the next round of evaluation has only a single candidate. In a call to `g(1, 1)`, the call is ambiguous even though `#2` has more `enable_if` attributes, because the first `enable_if` expressions are not ODR-equivalent.

Query for this feature with `__has_attribute(enable_if)`.

Note that functions with one or more `enable_if` attributes may not have their address taken, unless all of the conditions specified by said `enable_if` are constants that evaluate to `true`. For example:

```
const int TrueConstant = 1;
const int FalseConstant = 0;
int f(int a) __attribute__((enable_if(a > 0, "")));
int g(int a) __attribute__((enable_if(a == 0 | a != 0, "")));
int h(int a) __attribute__((enable_if(1, "")));
int i(int a) __attribute__((enable_if(TrueConstant, "")));
```

```
int j(int a) __attribute__((enable_if(FalseConstant, "")));

void fn() {
    int (*ptr)(int);
    ptr = &f; // error: 'a > 0' is not always true
    ptr = &g; // error: 'a == 0 || a != 0' is not a truthy constant
    ptr = &h; // OK: 1 is a truthy constant
    ptr = &i; // OK: 'TrueConstant' is a truthy constant
    ptr = &j; // error: 'FalseConstant' is a constant, but not truthy
}
```

Because `enable_if` evaluation happens during overload resolution, `enable_if` may give unintuitive results when used with templates, depending on when overloads are resolved. In the example below, clang will emit a diagnostic about no viable overloads for `foo` in `bar`, but not in `baz`:

```
double foo(int i) __attribute__((enable_if(i > 0, "")));
void *foo(int i) __attribute__((enable_if(i <= 0, "")));
template <int I>
auto bar() { return foo(I); }

template <typename T>
auto baz() { return foo(T::number); }

struct WithNumber { constexpr static int number = 1; };
void callThem() {
    bar<sizeof(WithNumber)>();
    baz<WithNumber>();
}
```

This is because, in `bar`, `foo` is resolved prior to template instantiation, so the value for `I` isn't known (thus, both `enable_if` conditions for `foo` fail). However, in `baz`, `foo` is resolved during template instantiation, so the value for `T::number` is known.

## external\_source\_symbol (clang::external\_source\_symbol)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
-----	-------	-----	------------	---------	--------	------------------------

X

X

X

The `external_source_symbol` attribute specifies that a declaration originates from an external source and describes the nature of that source.

The fact that Clang is capable of recognizing declarations that were defined externally can be used to provide better tooling support for mixed-language projects or projects that rely on auto-generated code. For instance, an IDE that uses Clang and that supports mixed-language projects can use this attribute to provide a correct ‘jump-to-definition’ feature. For a concrete example, consider a protocol that’s defined in a Swift file:

```
@objc public protocol SwiftProtocol {
    func method()
}
```

This protocol can be used from Objective-C code by including a header file that was generated by the Swift compiler. The declarations in that header can use the `external_source_symbol` attribute to make Clang aware of the fact that `SwiftProtocol` actually originates from a Swift module:

```
_attribute__((external_source_symbol(language="Swift",defined_in="module")))
@protocol SwiftProtocol
@required
- (void) method;
@end
```

Consequently, when ‘jump-to-definition’ is performed at a location that references `SwiftProtocol`, the IDE can jump to the original definition in the Swift source file rather than jumping to the Objective-C declaration in the auto-generated header file.

The `external_source_symbol` attribute is a comma-separated list that includes clauses that describe the origin and the nature of the particular declaration. Those clauses can be:

`language=string-literal`

The name of the source language in which this declaration was defined.

`defined_in=string-literal`

The name of the source container in which the declaration was defined. The exact definition of source container is language-specific, e.g. Swift’s source containers are modules, so `defined_in` should specify the Swift module name.

`generated_declaration`

This declaration was automatically generated by some tool.

The clauses can be specified in any order. The clauses that are listed above are all optional, but the attribute has to have at least one clause.

## flatten (gnu::flatten)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `flatten` attribute causes calls within the attributed function to be inlined unless it is impossible to do so, for example if the body of the callee is unavailable or if the callee has the `noinline` attribute.

## force\_align\_arg\_pointer (gnu::force\_align\_arg\_pointer)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Use this attribute to force stack alignment.

Legacy x86 code uses 4-byte stack alignment. Newer aligned SSE instructions (like ‘`movaps`’) that work with the stack require operands to be 16-byte aligned. This attribute realigns the stack in the function prologue to make sure the stack can be used with SSE instructions.

Note that the `x86_64` ABI forces 16-byte stack alignment at the call site. Because of this, ‘`force_align_arg_pointer`’ is not needed on `x86_64`, except in rare cases where the caller does not align the stack properly (e.g. flow jumps from `i386` arch code).

```
__attribute__((force_align_arg_pointer))
void f() {
    ...
}
```

## format (gnu::format)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Clang supports the `format` attribute, which indicates that the function accepts a `printf` or `scanf`-like format string and corresponding arguments or a `va_list` that contains these arguments.

Please see [GCC documentation about format attribute](#) to find details about attribute syntax.

Clang implements two kinds of checks with this attribute.

1. Clang checks that the function with the `format` attribute is called with a format string that uses format specifiers that are allowed, and that arguments match the format string. This is the `-Wformat` warning, it is on by default.
2. Clang checks that the format string argument is a literal string. This is the `-Wformat-nonliteral` warning, it is off by default.

Clang implements this mostly the same way as GCC, but there is a difference for functions that accept a `va_list` argument (for example, `vprintf`). GCC does not emit `-Wformat-nonliteral` warning for calls to such functions. Clang does not warn if the format string comes from a function parameter, where the function is annotated with a compatible attribute, otherwise it warns. For example:

```
__attribute__((__format__(__scanf__, 1, 3)))
void foo(const char* s, char *buf, ...) {
    va_list ap;
    va_start(ap, buf);

    vprintf(s, ap); // warning: format string is not a string literal
}
```

In this case we warn because `s` contains a format string for a `scanf`-like function, but it is passed to a `printf`-like function.

If the attribute is removed, clang still warns, because the format string is not a string literal.

Another example:

```
__attribute__((__format__(__printf__, 1, 3)))
void foo(const char* s, char *buf, ...) {
    va_list ap;
    va_start(ap, buf);
```

```
vprintf(s, ap); // warning
}
```

In this case Clang does not warn because the format string `s` and the corresponding arguments are annotated. If the arguments are incorrect, the caller of `foo` will receive a warning.

## ifunc (gnu::ifunc)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

`__attribute__((ifunc("resolver")))` is used to mark that the address of a declaration should be resolved at runtime by calling a resolver function.

The symbol name of the resolver function is given in quotes. A function with this name (after mangling) must be defined in the current translation unit; it may be `static`. The resolver function should take no arguments and return a pointer.

The `ifunc` attribute may only be used on a function declaration. A function declaration with an `ifunc` attribute is considered to be a definition of the declared entity. The entity must not have weak linkage; for example, in C++, it cannot be applied to a declaration if a definition at that location would be considered inline.

Not all targets support this attribute. ELF targets support this attribute when using `binutils v2.20.1` or higher and `glibc v2.11.1` or higher. Non-ELF targets currently do not support this attribute.

## internal\_linkage (clang::internal\_linkage)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `internal_linkage` attribute changes the linkage type of the declaration to internal. This is similar to C-style `static`, but can be used on classes and class methods. When applied to a class definition, this attribute affects all methods and static data members of that class. This can be used to contain the ABI of a C++ library by excluding unwanted class methods from the export tables.

## interrupt (ARM)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Clang supports the GNU style `__attribute__((interrupt("TYPE")))` attribute on ARM targets. This attribute may be attached to a function definition and instructs the backend to generate appropriate function entry/exit code so that it can be used directly as an interrupt service routine.

The parameter passed to the interrupt attribute is optional, but if provided it must be a string literal with one of the following values: “IRQ”, “FIQ”, “SWI”, “ABORT”, “UNDEF”.

The semantics are as follows:

If the function is AAPCS, Clang instructs the backend to realign the stack to 8 bytes on entry. This is a general requirement of the AAPCS at public interfaces, but may not hold when an exception is taken. Doing this allows other AAPCS functions to be called.

If the CPU is M-class this is all that needs to be done since the architecture itself is designed in such a way that functions obeying the normal AAPCS ABI constraints are valid exception handlers.

If the CPU is not M-class, the prologue and epilogue are modified to save all non-banked registers that are used, so that upon return the user-mode state will not be corrupted. Note that to avoid unnecessary overhead, only general-purpose (integer) registers are saved in this way. If VFP operations are needed, that state must be saved manually.

Specifically, interrupt kinds other than “FIQ” will save all core registers except “lr” and “sp”. “FIQ” interrupts will save r0-r7.

If the CPU is not M-class, the return instruction is changed to one of the canonical sequences permitted by the architecture for exception return. Where possible the function itself will make the necessary “lr” adjustments so that the “preferred return address” is selected.

Unfortunately the compiler is unable to make this guarantee for an “UNDEF” handler, where the offset from “lr” to the preferred return address depends on the execution state of the code which generated the exception. In this case a sequence equivalent to “movs pc, lr” will be used.

## interrupt (AVR)

### Supported Syntaxes



GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the GNU style `__attribute__((interrupt))` attribute on AVR targets. This attribute may be attached to a function definition and instructs the backend to generate appropriate function entry/exit code so that it can be used directly as an interrupt service routine.

On the AVR, the hardware globally disables interrupts when an interrupt is executed. The first instruction of an interrupt handler declared with this attribute is a SEI instruction to re-enable interrupts. See also the signal attribute that does not insert a SEI instruction.

## interrupt (MIPS)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the GNU style `__attribute__((interrupt("ARGUMENT")))` attribute on MIPS targets. This attribute may be attached to a function definition and instructs the backend to generate appropriate function entry/exit code so that it can be used directly as an interrupt service routine.

By default, the compiler will produce a function prologue and epilogue suitable for an interrupt service routine that handles an External Interrupt Controller (eic) generated interrupt. This behaviour can be explicitly requested with the “eic” argument.

Otherwise, for use with vectored interrupt mode, the argument passed should be of the form “vector=LEVEL” where LEVEL is one of the following values: “sw0”, “sw1”, “hw0”, “hw1”, “hw2”, “hw3”, “hw4”, “hw5”. The compiler will then set the interrupt mask to the corresponding level which will mask all interrupts up to and including the argument.

The semantics are as follows:

- The prologue is modified so that the Exception Program Counter (EPC) and Status coprocessor registers are saved to the stack. The interrupt mask is set so that the function can only be interrupted by a higher priority interrupt. The epilogue will restore the previous values of EPC and Status.

- The prologue and epilogue are modified to save and restore all non-kernel registers as necessary.

- The FPU is disabled in the prologue, as the floating pointer registers are not spilled to the stack.

- The function return sequence is changed to use an exception return instruction.

The parameter sets the interrupt mask for the function corresponding to the interrupt level specified. If no mask is specified the interrupt mask defaults to “eic”.

## kernel

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

`__attribute__((kernel))` is used to mark a kernel function in RenderScript.

In RenderScript, kernel functions are used to express data-parallel computations. The RenderScript runtime efficiently parallelizes kernel functions to run on computational resources such as multi-core CPUs and GPUs. See the **RenderScript** documentation for more information.

## long\_call (gnu::long\_call, gnu::far)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the `__attribute__((long_call))`, `__attribute__((far))`, and `__attribute__((near))` attributes on MIPS targets. These attributes may only be added to function declarations and change the code generated by the compiler when directly calling the function. The `near` attribute allows calls to the function to be made using the `jal` instruction, which requires the function to be located in the same naturally aligned 256MB segment as the caller. The `long_call` and `far` attributes are synonyms and require the use of a different call sequence that works regardless of the distance between the functions.

These attributes have no effect for position-independent code.

These attributes take priority over command line switches such as `-mlong-calls` and `-mno-long-calls`.

## micromips (gnu::micromips)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the GNU style `__attribute__((micromips))` and `__attribute__((nommicromips))` attributes on MIPS targets. These attributes may be attached to a function definition and instructs the backend to generate or not to generate microMIPS code for that function.

These attributes override the `-mmicromips` and `-mno-micromips` options on the command line.

### **no\_caller\_saved\_registers (gnu::no\_caller\_saved\_registers)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Use this attribute to indicate that the specified function has no caller-saved registers. That is, all registers are callee-saved except for registers used for passing parameters to the function or returning parameters from the function. The compiler saves and restores any modified registers that were not used for passing or returning arguments to the function.

The user can call functions specified with the ‘no\_caller\_saved\_registers’ attribute from an interrupt handler without saving and restoring all call-clobbered registers.

Note that ‘no\_caller\_saved\_registers’ attribute is not a calling convention. In fact, it only overrides the decision of which registers should be saved by the caller, but not how the parameters are passed from the caller to the callee.

For example:

```
__attribute__((no_caller_saved_registers, fastcall))
void f(int arg1, int arg2) {
    ...
}
```

In this case parameters 'arg1' and 'arg2' will be passed in registers. In this case, on 32-bit x86 targets, the function 'f' will use ECX and EDX as register parameters. However, it will not assume any scratch registers and should save and restore any modified registers except for ECX and EDX.

## no\_sanitiz (clang::no\_sanitiz)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Use the `no_sanitiz` attribute on a function declaration to specify that a particular instrumentation or set of instrumentations should not be applied to that function. The attribute takes a list of string literals, which have the same meaning as values accepted by the `-fno-sanitiz=` flag. For example, `__attribute__((no_sanitiz("address", "thread")))` specifies that AddressSanitizer and ThreadSanitizer should not be applied to the function.

See **Controlling Code Generation** for a full list of supported sanitizer flags.

## no\_sanitiz\_address (no\_address\_safety\_analysis, gnu::no\_address\_safety\_analysis, gnu::no\_sanitiz\_address)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Use `__attribute__((no_sanitiz_address))` on a function declaration to specify that address safety instrumentation (e.g. AddressSanitizer) should not be applied to that function.

## no\_sanitiz\_memory

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Use `__attribute__((no_sanitize_memory))` on a function declaration to specify that checks for uninitialized memory should not be inserted (e.g. by MemorySanitizer). The function may still be instrumented by the tool to avoid false positives in other places.

## no\_sanitize\_thread

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Use `__attribute__((no_sanitize_thread))` on a function declaration to specify that checks for data races on plain (non-atomic) memory accesses should not be inserted by ThreadSanitizer. The function is still instrumented by the tool to avoid false positives and provide meaningful stack traces.

## no\_split\_stack (gnu::no\_split\_stack)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `no_split_stack` attribute disables the emission of the split stack preamble for a particular function. It has no effect if `-fsplit-stack` is not specified.

## noalias

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
			X			

The `noalias` attribute indicates that the only memory accesses inside function are loads and stores from objects pointed to by its pointer-typed arguments, with arbitrary offsets.

**nodiscard, warn\_unused\_result, clang::warn\_unused\_result, gnu::warn\_unused\_result**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X	X				X

Clang supports the ability to diagnose when the results of a function call expression are discarded under suspicious circumstances. A diagnostic is generated when a function or its return type is marked with `[[nodiscard]]` (or `__attribute__((warn_unused_result))`) and the function call appears as a potentially-evaluated discarded-value expression that is not explicitly cast to `void`.

**noduplicate (clang::noduplicate)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `noduplicate` attribute can be placed on function declarations to control whether function calls to this function can be duplicated or not as a result of optimizations. This is required for the implementation of functions with certain special requirements, like the OpenCL “barrier” function, that might need to be run concurrently by all the threads that are executing in lockstep on the hardware. For example this attribute applied on the function “`nodupfunc`” in the code below avoids that:

```
void nodupfunc() __attribute__((noduplicate));
// Setting it as a C++11 attribute is also valid
// void nodupfunc() [[clang::noduplicate]];
void foo();
void bar();

nodupfunc();
if (a > n) {
    foo();
} else {
    bar();
}
```

gets possibly modified by some optimizations into code similar to this:

```
if (a > n) {
    nodupfunc();
    foo();
} else {
    nodupfunc();
    bar();
}
```

where the call to “nodupfunc” is duplicated and sunk into the two branches of the condition.

## **nomicromips (gnu::nomicromips)**

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the GNU style `__attribute__((micromips))` and `__attribute__((nomicromips))` attributes on MIPS targets. These attributes may be attached to a function definition and instructs the backend to generate or not to generate microMIPS code for that function.

These attributes override the `-mmicromips` and `-mno-micromips` options on the command line.

## **noreturn**

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
	X					X

A function declared as `[[noreturn]]` shall not return to its caller. The compiler will generate a diagnostic for a function declared as `[[noreturn]]` that appears to be capable of returning to its caller.

## not\_tail\_called (clang::not\_tail\_called)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `not_tail_called` attribute prevents tail-call optimization on statically bound calls. It has no effect on indirect calls. Virtual functions, objective-c methods, and functions marked as `always_inline` cannot be marked as `not_tail_called`.

For example, it prevents tail-call optimization in the following case:

```
int __attribute__((not_tail_called)) foo1(int);

int foo2(int a) {
    return foo1(a); // No tail-call optimization on direct calls.
}
```

However, it doesn't prevent tail-call optimization in this case:

```
int __attribute__((not_tail_called)) foo1(int);

int foo2(int a) {
    int (*fn)(int) = &foo1;

    // not_tail_called has no effect on an indirect call even if the call can be
    // resolved at compile time.
    return (*fn)(a);
}
```

Marking virtual functions as `not_tail_called` is an error:

```
class Base {
public:
    // not_tail_called on a virtual function is an error.
    [[clang::not_tail_called]] virtual int foo1();
}
```



```

virtual int foo2();

// Non-virtual functions can be marked ``not_tail_called``.
[[clang::not_tail_called]] int foo3();
};

class Derived1 : public Base {
public:
    int foo1() override;

    // not_tail_called on a virtual function is an error.
    [[clang::not_tail_called]] int foo2() override;
};

```

## nothrow (gnu::nothrow)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X		X			X

Clang supports the GNU style `_attribute__((nothrow))` and Microsoft style `__declspec(nothrow)` attribute as an equivalent of *noexcept* on function declarations. This attribute informs the compiler that the annotated function does not throw an exception. This prevents exception-unwinding. This attribute is particularly useful on functions in the C Standard Library that are guaranteed to not throw an exception.

## objc\_boxable

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Structs and unions marked with the `objc_boxable` attribute can be used with the Objective-C boxed expression syntax, `@(...)`.

**Usage:** `__attribute__((objc_boxable))`. This attribute can only be placed on a declaration of a trivially-copyable struct or union:

```
struct __attribute__((objc_boxable)) some_struct {
    int i;
};
union __attribute__((objc_boxable)) some_union {
    int i;
    float f;
};
typedef struct __attribute__((objc_boxable)) _some_struct some_struct;

// ...

some_struct ss;
NSValue *boxed = @(ss);
```

## objc\_method\_family

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Many methods in Objective-C have conventional meanings determined by their selectors. It is sometimes useful to be able to mark a method as having a particular conventional meaning despite not having the right selector, or as not having the conventional meaning that its selector would suggest. For these use cases, we provide an attribute to specifically describe the “method family” that a method belongs to.

**Usage:** `__attribute__((objc_method_family(X)))`, where X is one of `none`, `alloc`, `copy`, `init`, `mutableCopy`, or `new`. This attribute can only be placed at the end of a method declaration:

```
-(NSString *)initMyStringValue __attribute__((objc_method_family(none)));
```

Users who do not wish to change the conventional meaning of a method, and who merely want to document its non-standard retain and release semantics, should use the retaining behavior attributes (`ns_returns_retained`, `ns_returns_not_retained`, etc).

Query for this feature with `__has_attribute(objc_method_family)`.

## objc\_requires\_super

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Some Objective-C classes allow a subclass to override a particular method in a parent class but expect that the overriding method also calls the overridden method in the parent class. For these cases, we provide an attribute to designate that a method requires a “call to super” in the overriding method in the subclass.

**Usage:** `__attribute__((objc_requires_super))`. This attribute can only be placed at the end of a method declaration:

```
- (void)foo __attribute__((objc_requires_super));
```

This attribute can only be applied the method declarations within a class, and not a protocol. Currently this attribute does not enforce any placement of where the call occurs in the overriding method (such as in the case of `-dealloc` where the call must appear at the end). It checks only that it exists.

Note that on both OS X and iOS that the Foundation framework provides a convenience macro `NS_REQUIRES_SUPER` that provides syntactic sugar for this attribute:

```
- (void)foo NS_REQUIRES_SUPER;
```

This macro is conditionally defined depending on the compiler’s support for this attribute. If the compiler does not support the attribute the macro expands to nothing.

Operationally, when a method has this annotation the compiler will warn if the implementation of an override in a subclass does not call super. For example:

```
warning: method possibly missing a [super AnnotMeth] call
- (void) AnnotMeth{};
```

## objc\_runtime\_name

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

By default, the Objective-C interface or protocol identifier is used in the metadata name for that object. The *objc\_runtime\_name* attribute allows annotated interfaces or protocols to use the specified string argument in the object's metadata name instead of the default name.

**Usage:** `__attribute__((objc_runtime_name("MyLocalName")))`. This attribute can only be placed before an `@protocol` or `@interface` declaration:

```
__attribute__((objc_runtime_name("MyLocalName")))
@interface Message
@end
```

## objc\_runtime\_visible

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

This attribute specifies that the Objective-C class to which it applies is visible to the Objective-C runtime but not to the linker. Classes annotated with this attribute cannot be subclassed and cannot have categories defined for them.

## optnone (clang::optnone)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `optnone` attribute suppresses essentially all optimizations on a function or method, regardless of the optimization level applied to the compilation unit as a whole. This is particularly useful when you need to debug a particular function, but it is infeasible to build the entire application without optimization. Avoiding optimization on the specified function can improve the quality of the debugging information for that function.

This attribute is incompatible with the `always_inline` and `minsize` attributes.

## overloadable

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Clang provides support for C++ function overloading in C. Function overloading in C is introduced using the `overloadable` attribute. For example, one might provide several overloaded versions of a `tgysin` function that invokes the appropriate standard function computing the sine of a value with float, double, or long double precision:

```
#include <math.h>
float __attribute__((overloadable)) tgysin(float x) { return sinf(x); }
double __attribute__((overloadable)) tgysin(double x) { return sin(x); }
long double __attribute__((overloadable)) tgysin(long double x) { return sinl(x); }
```

Given these declarations, one can call `tgysin` with a float value to receive a float result, with a double to receive a double result, etc. Function overloading in C follows the rules of C++ function overloading to pick the best overload given the call arguments, with a few C-specific semantics:

Conversion from float or double to long double is ranked as a floating-point promotion (per C99) rather than as a floating-point conversion (as in C++).

A conversion from a pointer of type `T*` to a pointer of type `U*` is considered a pointer conversion (with conversion rank) if `T` and `U` are compatible types.

A conversion from type  $\tau$  to a value of type  $u$  is permitted if  $\tau$  and  $u$  are compatible types. This conversion is given “conversion” rank. If no viable candidates are otherwise available, we allow a conversion from a pointer of type  $\tau^*$  to a pointer of type  $u^*$ , where  $\tau$  and  $u$  are incompatible. This conversion is ranked below all other types of conversions. Please note:  $u$  lacking qualifiers that are present on  $\tau$  is sufficient for  $\tau$  and  $u$  to be incompatible.

The declaration of overloadable functions is restricted to function declarations and definitions. If a function is marked with the `overloadable` attribute, then all declarations and definitions of functions with that name, except for at most one (see the note below about unmarked overloads), must have the `overloadable` attribute. In addition, redeclarations of a function with the `overloadable` attribute must have the `overloadable` attribute, and redeclarations of a function without the `overloadable` attribute must *not* have the `overloadable` attribute. e.g.,

```
int f(int) __attribute__((overloadable));
float f(float); // error: declaration of "f" must have the "overloadable" attribute
int f(int); // error: redeclaration of "f" must have the "overloadable" attribute

int g(int) __attribute__((overloadable));
int g(int) {} // error: redeclaration of "g" must also have the "overloadable" attribute

int h(int);
int h(int) __attribute__((overloadable)); // error: declaration of "h" must not
// have the "overloadable" attribute
```

Functions marked `overloadable` must have prototypes. Therefore, the following code is ill-formed:

```
int h() __attribute__((overloadable)); // error: h does not have a prototype
```

However, `overloadable` functions are allowed to use a ellipsis even if there are no named parameters (as is permitted in C++). This feature is particularly useful when combined with the `unavailable` attribute:

```
void honeypot(...) __attribute__((overloadable, unavailable)); // calling me is an error
```

Functions declared with the `overloadable` attribute have their names mangled according to the same rules as C++ function names. For example, the three `tgssin` functions in our motivating example get the mangled names `_Z5tgssinf`, `_Z5tgssind`, and `_Z5tgssine`, respectively. There are two caveats to this use of name mangling:

Future versions of Clang may change the name mangling of functions overloaded in C, so you should not depend on an specific mangling. To be completely safe, we strongly urge the use of `static inline` with `overloadable` functions.

The `overloadable` attribute has almost no meaning when used in C++, because names will already be mangled and functions are already overloadable. However, when an `overloadable` function occurs within an `extern "C"` linkage specification, it's name *will* be mangled in the same way as it would in C.

For the purpose of backwards compatibility, at most one function with the same name as other `overloadable` functions may omit the `overloadable` attribute. In this case, the function without the `overloadable` attribute will not have its name mangled.

For example:

```
// Notes with mangled names assume Itanium mangling.
int f(int);
int f(double) __attribute__((overloadable));
void foo() {
    f(5); // Emits a call to f (not _Z1fi, as it would with an overload that
         // was marked with overloadable).
    f(1.0); // Emits a call to _Z1fd.
}
```

Support for unmarked overloads is not present in some versions of clang. You may query for it using `__has_extension(overloadable_unmarked)`.

Query for this attribute with `__has_attribute(overloadable)`.

## **release\_capability (release\_shared\_capability, clang::release\_capability, clang::release\_shared\_capability)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Marks a function as releasing a capability.

## **short\_call (gnu::short\_call, gnu::near)**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the `__attribute__((long_call))`, `__attribute__((far))`, `__attribute__((short_call))`, and `__attribute__((near))` attributes on MIPS targets. These attributes may only be added to function declarations and change the code generated by the compiler when directly calling the function. The `short_call` and `near` attributes are synonyms and allow calls to the function to be made using the `jal` instruction, which requires the function to be located in the same naturally aligned 256MB segment as the caller. The `long_call` and `far` attributes are synonyms and require the use of a different call sequence that works regardless of the distance between the functions.

These attributes have no effect for position-independent code.

These attributes take priority over command line switches such as `-mlong-calls` and `-mno-long-calls`.

## signal (gnu::signal)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Clang supports the GNU style `__attribute__((signal))` attribute on AVR targets. This attribute may be attached to a function definition and instructs the backend to generate appropriate function entry/exit code so that it can be used directly as an interrupt service routine.

Interrupt handler functions defined with the `signal` attribute do not re-enable interrupts.

## target (gnu::target)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X



Clang supports the GNU style `__attribute__((target("OPTIONS")))` attribute. This attribute may be attached to a function definition and instructs the backend to use different code generation options than were passed on the command line.

The current set of options correspond to the existing “subtarget features” for the target with or without a “-mno-” in front corresponding to the absence of the feature, as well as `arch="CPU"` which will change the default “CPU” for the function.

Example “subtarget features” from the x86 backend include: “mmx”, “sse”, “sse4.2”, “avx”, “xop” and largely correspond to the machine specific options handled by the front end.

**`try_acquire_capability (try_acquire_shared_capability, clang::try_acquire_capability, clang::try_acquire_shared_capability)`**

#### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

Marks a function that attempts to acquire a capability. This function may fail to actually acquire the capability; they accept a Boolean value determining whether acquiring the capability means success (true), or failing to acquire the capability means success (false).

**`xray_always_instrument (clang::xray_always_instrument), xray_never_instrument (clang::xray_never_instrument), xray_log_args (clang::xray_log_args)`**

#### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

`__attribute__((xray_always_instrument))` or `[[clang::xray_always_instrument]]` is used to mark member functions (in C++), methods (in Objective C), and free functions (in C, C++, and Objective C) to be instrumented with XRay. This will cause the function to always have space at the beginning and exit points to allow for runtime patching.

Conversely, `__attribute__((xray_never_instrument))` or `[[clang::xray_never_instrument]]` will inhibit the insertion of these instrumentation points.

If a function has neither of these attributes, they become subject to the XRay heuristics used to determine whether a function should be instrumented or otherwise.

`__attribute__((xray_log_args(N)))` or `[[clang::xray_log_args(N)]]` is used to preserve N function arguments for the logging function. Currently, only N==1 is supported.

**`xray_always_instrument (clang::xray_always_instrument)`, `xray_never_instrument (clang::xray_never_instrument)`, `xray_log_args (clang::xray_log_args)`**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

`__attribute__((xray_always_instrument))` or `[[clang::xray_always_instrument]]` is used to mark member functions (in C++), methods (in Objective C), and free functions (in C, C++, and Objective C) to be instrumented with XRay. This will cause the function to always have space at the beginning and exit points to allow for runtime patching.

Conversely, `__attribute__((xray_never_instrument))` or `[[clang::xray_never_instrument]]` will inhibit the insertion of these instrumentation points.

If a function has neither of these attributes, they become subject to the XRay heuristics used to determine whether a function should be instrumented or otherwise.

`__attribute__((xray_log_args(N)))` or `[[clang::xray_log_args(N)]]` is used to preserve N function arguments for the logging function. Currently, only N==1 is supported.

## Variable Attributes

**`dllexport (gnu::dllexport)`**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute

X	X	X	X
---	---	---	---

The `__declspec(dllexport)` attribute declares a variable, function, or Objective-C interface to be exported from the module. It is available under the `-fdeclspec` flag for compatibility with various compilers. The primary use is for COFF object files which explicitly specify what interfaces are available for external use. See the [dllexport](#) documentation on MSDN for more information.

## **dllimport (gnu::dllimport)**

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X		X			X

The `__declspec(dllimport)` attribute declares a variable, function, or Objective-C interface to be imported from an external module. It is available under the `-fdeclspec` flag for compatibility with various compilers. The primary use is for COFF object files which explicitly specify what interfaces are imported from external modules. See the [dllimport](#) documentation on MSDN for more information.

## **init\_seg**

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
					X	

The attribute applied by `pragma init_seg()` controls the section into which global initialization function pointers are emitted. It is only available with `-fms-extensions`. Typically, this function pointer is emitted into `.CRT$XCU` on Windows. The user can change the order of initialization by using a different section name with the same `.CRT$XC` prefix and a suffix that sorts lexicographically before or after the standard `.CRT$XCU` sections. See the [init\\_seg](#) documentation on MSDN for more information.

## **maybe\_unused, unused, gnu::unused**

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X	X				

When passing the `-Wunused` flag to Clang, entities that are unused by the program may be diagnosed. The `[[maybe_unused]]` (or `__attribute__((unused))`) attribute can be used to silence such diagnostics when the entity cannot be removed. For instance, a local variable may exist solely for use in an `assert()` statement, which makes the local variable unused when `NDEBUG` is defined.

The attribute may be applied to the declaration of a class, a typedef, a variable, a function or method, a function parameter, an enumeration, an enumerator, a non-static data member, or a label.

### nodebug (gnu::nodebug)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `nodebug` attribute allows you to suppress debugging information for a function or method, or for a variable that is not a parameter or a non-static data member.

### noescape (clang::noescape)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

`noescape` placed on a function parameter of a pointer type is used to inform the compiler that the pointer cannot escape: that is, no reference to the object the pointer points to that is derived from the parameter value will survive after the function returns. Users are responsible for making sure parameters annotated with `noescape` do not actually escape.

For example:

```
int *gp;

void nonescapingFunc(__attribute__((noescape)) int *p) {
    *p += 100; // OK.
}

void escapingFunc(__attribute__((noescape)) int *p) {
    gp = p; // Not OK.
}
```

Additionally, when the parameter is a *block pointer* <<https://clang.llvm.org/docs/BlockLanguageSpec.html>>, the same restriction applies to copies of the block. For example:

```
typedef void (^BlockTy)();
BlockTy g0, g1;

void nonescapingFunc(__attribute__((noescape)) BlockTy block) {
    block(); // OK.
}

void escapingFunc(__attribute__((noescape)) BlockTy block) {
    g0 = block; // Not OK.
    g1 = Block_copy(block); // Not OK either.
}
```

## nosvm

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

OpenCL 2.0 supports the optional `__attribute__((nosvm))` qualifier for pointer variable. It informs the compiler that the pointer does not refer to a shared virtual memory region. See OpenCL v2.0 s6.7.2 for details.

Since it is not widely used and has been removed from OpenCL 2.1, it is ignored by Clang.

## pass\_object\_size

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

#### Note

The mangling of functions with parameters that are annotated with `pass_object_size` is subject to change. You can get around this by using `__asm__("foo")` to explicitly name your functions, thus preserving your ABI; also, non-overloadable C functions with `pass_object_size` are not mangled.

The `pass_object_size(Type)` attribute can be placed on function parameters to instruct clang to call `__builtin_object_size(param, Type)` at each callsite of said function, and implicitly pass the result of this call in as an invisible argument of type `size_t` directly after the parameter annotated with `pass_object_size`. Clang will also replace any calls to `__builtin_object_size(param, Type)` in the function by said implicit parameter.

Example usage:

```
int bzero1(char *const p __attribute__((pass_object_size(0))))
__attribute__((noinline)) {
    int i = 0;
    for (/**/; i < (int)__builtin_object_size(p, 0); ++i) {
        p[i] = 0;
    }
    return i;
}

int main() {
    char chars[100];
    int n = bzero1(&chars[0]);
    assert(n == sizeof(chars));
}
```

```
    return 0;
}
```

If successfully evaluating `__builtin_object_size(param, Type)` at the callsite is not possible, then the “failed” value is passed in. So, using the definition of `bzero1` from above, the following code would exit cleanly:

```
int main2(int argc, char *argv[]) {
    int n = bzero1(argv);
    assert(n == -1);
    return 0;
}
```

`pass_object_size` plays a part in overload resolution. If two overload candidates are otherwise equally good, then the overload with one or more parameters with `pass_object_size` is preferred. This implies that the choice between two identical overloads both with `pass_object_size` on one or more parameters will always be ambiguous; for this reason, having two such overloads is illegal. For example:

```
#define PS(N) __attribute__((pass_object_size(N)))
// OK
void Foo(char *a, char *b); // Overload A
// OK -- overload A has no parameters with pass_object_size.
void Foo(char *a PS(0), char *b PS(0)); // Overload B
// Error -- Same signature (sans pass_object_size) as overload B, and both
// overloads have one or more parameters with the pass_object_size attribute.
void Foo(void *a PS(0), void *b);

// OK
void Bar(void *a PS(0)); // Overload C
// OK
void Bar(char *c PS(1)); // Overload D

void main() {
    char known[10], *unknown;
    Foo(unknown, unknown); // Calls overload B
    Foo(known, unknown); // Calls overload B
    Foo(unknown, known); // Calls overload B
    Foo(known, known); // Calls overload B
```

```
Bar(known); // Calls overload D
Bar(unknown); // Calls overload D
}
```

Currently, `pass_object_size` is a bit restricted in terms of its usage:

Only one use of `pass_object_size` is allowed per parameter.

It is an error to take the address of a function with `pass_object_size` on any of its parameters. If you wish to do this, you can create an overload without `pass_object_size` on any parameters.

It is an error to apply the `pass_object_size` attribute to parameters that are not pointers. Additionally, any parameter that `pass_object_size` is applied to must be marked `const` at its function's definition.

## require\_constant\_initialization (clang::require\_constant\_initialization)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

This attribute specifies that the variable to which it is attached is intended to have a **constant initializer** according to the rules of [basic.start.static]. The variable is required to have static or thread storage duration. If the initialization of the variable is not a constant initializer an error will be produced. This attribute may only be used in C++.

Note that in C++03 strict constant expression checking is not done. Instead the attribute reports if Clang can emit the variable as a constant, even if it's not technically a 'constant initializer'. This behavior is non-portable.

Static storage duration variables with constant initializers avoid hard-to-find bugs caused by the indeterminate order of dynamic initialization. They can also be safely used during dynamic initialization across translation units.

This attribute acts as a compile time assertion that the requirements for constant initialization have been met. Since these requirements change between dialects and have subtle pitfalls it's important to fail fast instead of silently falling back on dynamic initialization.

```
// -std=c++14
#define SAFE_STATIC [[clang::require_constant_initialization]]
struct T {
    constexpr T(int) {}
}
```



```

~T(); // non-trivial
};
SAFE_STATIC T x = {42}; // Initialization OK. Doesn't check destructor.
SAFE_STATIC T y = 42; // error: variable does not have a constant initializer
// copy initialization is not a constant expression on a non-literal type.

```

## section (gnu::section, \_\_declspec(allocate))

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X		X			X

The `section` attribute allows you to specify a specific section a global variable or function should be in after translation.

## swift\_context

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `swift_context` attribute marks a parameter of a `swiftcall` function as having the special context-parameter ABI treatment.

This treatment generally passes the context value in a special register which is normally callee-preserved.

A `swift_context` parameter must either be the last parameter or must be followed by a `swift_error_result` parameter (which itself must always be the last parameter).

A context parameter must have pointer or reference type.

## swift\_error\_result

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `swift_error_result` attribute marks a parameter of a `swiftcall` function as having the special error-result ABI treatment.

This treatment generally passes the underlying error value in and out of the function through a special register which is normally callee-preserved. This is modeled in C by pretending that the register is addressable memory:

The caller appears to pass the address of a variable of pointer type. The current value of this variable is copied into the register before the call; if the call returns normally, the value is copied back into the variable.

The callee appears to receive the address of a variable. This address is actually a hidden location in its own stack, initialized with the value of the register upon entry. When the function returns normally, the value in that hidden location is written back to the register.

A `swift_error_result` parameter must be the last parameter, and it must be preceded by a `swift_context` parameter.

A `swift_error_result` parameter must have type `T**` or `T*&` for some type `T`. Note that no qualifiers are permitted on the intermediate level.

It is undefined behavior if the caller does not pass a pointer or reference to a valid object.

The standard convention is that the error value itself (that is, the value stored in the apparent argument) will be null upon function entry, but this is not enforced by the ABI.

## swift\_indirect\_result

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `swift_indirect_result` attribute marks a parameter of a `swiftcall` function as having the special indirect-result ABI treatment.

This treatment gives the parameter the target's normal indirect-result ABI treatment, which may involve passing it differently from an ordinary parameter. However, only the first indirect result will receive this treatment. Furthermore, low-level lowering may decide that a direct result must be returned indirectly; if so, this will take priority over the `swift_indirect_result` parameters.

A `swift_indirect_result` parameter must either be the first parameter or follow another `swift_indirect_result` parameter.

A `swift_indirect_result` parameter must have type  $\tau^*$  or  $\tau\&$  for some object type  $\tau$ . If  $\tau$  is a complete type at the point of definition of a function, it is undefined behavior if the argument value does not point to storage of adequate size and alignment for a value of type  $\tau$ .

Making indirect results explicit in the signature allows C functions to directly construct objects into them without relying on language optimizations like C++’s named return value optimization (NRVO).

## swiftcall

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

The `swiftcall` attribute indicates that a function should be called using the Swift calling convention for a function or function pointer.

The lowering for the Swift calling convention, as described by the Swift ABI documentation, occurs in multiple phases. The first, “high-level” phase breaks down the formal parameters and results into innately direct and indirect components, adds implicit parameters for the generic signature, and assigns the context and error ABI treatments to parameters where applicable. The second phase breaks down the direct parameters and results from the first phase and assigns them to registers or the stack. The `swiftcall` convention only handles this second phase of lowering; the C function type must accurately reflect the results of the first phase, as follows:

Results classified as indirect by high-level lowering should be represented as parameters with the `swift_indirect_result` attribute.

Results classified as direct by high-level lowering should be represented as follows:

First, remove any empty direct results.

If there are no direct results, the C result type should be `void`.

If there is one direct result, the C result type should be a type with the exact layout of that result type.

If there are a multiple direct results, the C result type should be a struct type with the exact layout of a tuple of those results.

Parameters classified as indirect by high-level lowering should be represented as parameters of pointer type.

Parameters classified as direct by high-level lowering should be omitted if they are empty types; otherwise, they should be represented as a parameter type with a layout exactly matching the layout of the Swift parameter type.

The context parameter, if present, should be represented as a trailing parameter with the `swift_context` attribute.

The error result parameter, if present, should be represented as a trailing parameter (always following a context parameter) with the `swift_error_result` attribute.

`swiftcall` does not support variadic arguments or unprototyped functions.

The parameter ABI treatment attributes are aspects of the function type. A function type which applies an ABI treatment attribute to a parameter is a different type from an otherwise-identical function type that does not. A single parameter may not have multiple ABI treatment attributes.

Support for this feature is target-dependent, although it should be supported on every target that Swift supports. Query for this support with `__has_attribute(swiftcall)`. This implies support for the `swift_context`, `swift_error_result`, and `swift_indirect_result` attributes.

## thread

### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
			X			

The `__declspec(thread)` attribute declares a variable with thread local storage. It is available under the `-fms-extensions` flag for MSVC compatibility. See the documentation for **`__declspec(thread)`** on MSDN.

In Clang, `__declspec(thread)` is generally equivalent in functionality to the GNU `__thread` keyword. The variable must not have a destructor and must have a constant initializer, if any. The attribute only applies to variables declared with static storage duration, such as globals, class static data members, and static locals.

## tls\_model (gnu::tls\_model)

### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
X	X					X

The `tls_model` attribute allows you to specify which thread-local storage model to use. It accepts the following strings:

```
global-dynamic
local-dynamic
initial-exec
```

local-exec

TLS models are mutually exclusive.

## Type Attributes

### `__single_inheritance`, `__multiple_inheritance`, `__virtual_inheritance`

Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
				X		

This collection of keywords is enabled under `-fms-extensions` and controls the pointer-to-member representation used on `*-*-win32` targets.

The `*-*-win32` targets utilize a pointer-to-member representation which varies in size and alignment depending on the definition of the underlying class.

However, this is problematic when a forward declaration is only available and no definition has been made yet. In such cases, Clang is forced to utilize the most general representation that is available to it.

These keywords make it possible to use a pointer-to-member representation other than the most general one regardless of whether or not the definition will ever be present in the current translation unit.

This family of keywords belong between the `class-key` and `class-name`:

```
struct __single_inheritance S;
int S::*i;
struct S {};
```

This keyword can be applied to class templates but only has an effect when used on full specializations:

```
template <typename T, typename U> struct __single_inheritance A; // warning: inheritance model ignored on primary template
template <typename T> struct __multiple_inheritance A<T, T>; // warning: inheritance model ignored on partial specialization
template <> struct __single_inheritance A<int, float>;
```

Note that choosing an inheritance model less general than strictly necessary is an error:

```
struct __multiple_inheritance S; // error: inheritance model does not match definition
int S::*i;
struct S {};
```

## align\_value

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `align_value` attribute can be added to the typedef of a pointer type or the declaration of a variable of pointer or reference type. It specifies that the pointer will point to, or the reference will bind to, only objects with at least the provided alignment. This alignment value must be some positive power of 2.

```
typedef double * aligned_double_ptr __attribute__((align_value(64)));
void foo(double & x __attribute__((align_value(128)),
      aligned_double_ptr y) { ... }
```

If the pointer value does not have the specified alignment at runtime, the behavior of the program is undefined.

## empty\_bases

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
			X			

The `empty_bases` attribute permits the compiler to utilize the empty-base-optimization more frequently. This attribute only applies to struct, class, and union types. It is only supported when using the Microsoft C++ ABI.

**enum\_extensibility (clang::enum\_extensibility)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

Attribute `enum_extensibility` is used to distinguish between enum definitions that are extensible and those that are not. The attribute can take either `closed` or `open` as an argument. `closed` indicates a variable of the enum type takes a value that corresponds to one of the enumerators listed in the enum definition or, when the enum is annotated with `flag_enum`, a value that can be constructed using values corresponding to the enumerators. `open` indicates a variable of the enum type can take any values allowed by the standard and instructs clang to be more lenient when issuing warnings.

```
enum __attribute__((enum_extensibility(closed))) ClosedEnum {
    A0, A1
};

enum __attribute__((enum_extensibility(open))) OpenEnum {
    B0, B1
};

enum __attribute__((enum_extensibility(closed),flag_enum)) ClosedFlagEnum {
    C0 = 1 << 0, C1 = 1 << 1
};

enum __attribute__((enum_extensibility(open),flag_enum)) OpenFlagEnum {
    D0 = 1 << 0, D1 = 1 << 1
};

void foo1() {
    enum ClosedEnum ce;
    enum OpenEnum oe;
    enum ClosedFlagEnum cfe;
    enum OpenFlagEnum ofe;

    ce = A1;      // no warnings
    ce = 100;     // warning issued
}
```

```

oe = B1;      // no warnings
oe = 100;     // no warnings
cfe = C0 | C1; // no warnings
cfe = C0 | C1 | 4; // warning issued
ofe = D0 | D1; // no warnings
ofe = D0 | D1 | 4; // no warnings
}

```

## flag\_enum

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

This attribute can be added to an enumerator to signal to the compiler that it is intended to be used as a flag type. This will cause the compiler to assume that the range of the type includes all of the values that you can get by manipulating bits of the enumerator when issuing warnings.

## layout\_version

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
			X			

The layout\_version attribute requests that the compiler utilize the class layout rules of a particular compiler version. This attribute only applies to struct, class, and union types. It is only supported when using the Microsoft C++ ABI.

## lto\_visibility\_public (clang::lto\_visibility\_public)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute



X

X

See [LTO Visibility](#).

## novtable

Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
			X			

This attribute can be added to a class declaration or definition to signal to the compiler that constructors and destructors will not reference the virtual function table. It is only supported when using the Microsoft C++ ABI.

## objc\_subclassing\_restricted

Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
X						X

This attribute can be added to an Objective-C `@interface` declaration to ensure that this class cannot be subclassed.

## selectany (gnu::selectany)

Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
X	X		X			

This attribute appertains to a global symbol, causing it to have a weak definition ( [linkonce](#) ), allowing the linker to select any definition.

For more information see [gcc documentation](#) or [msvc documentation](#).

## transparent\_union (gnu::transparent\_union)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

This attribute can be applied to a union to change the behaviour of calls to functions that have an argument with a transparent union type. The compiler behaviour is changed in the following manner:

A value whose type is any member of the transparent union can be passed as an argument without the need to cast that value.

The argument is passed to the function using the calling convention of the first member of the transparent union. Consequently, all the members of the transparent union should have the same calling convention as its first member.

Transparent unions are not supported in C++.

## Statement Attributes

### #pragma clang loop

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
					X	

The `#pragma clang loop` directive allows loop optimization hints to be specified for the subsequent loop. The directive allows vectorization, interleaving, and unrolling to be enabled or disabled. Vector width as well as interleave and unrolling count can be manually specified. See [language extensions](#) for details.

### #pragma unroll, #pragma nounroll

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
					X	

Loop unrolling optimization hints can be specified with `#pragma unroll` and `#pragma nounroll`. The pragma is placed immediately before a for, while, do-while, or c++11 range-based for loop.

Specifying `#pragma unroll` without a parameter directs the loop unroller to attempt to fully unroll the loop if the trip count is known at compile time and attempt to partially unroll the loop if the trip count is not known at compile time:

```
#pragma unroll
for (...) {
    ...
}
```

Specifying the optional parameter, `#pragma unroll _value_`, directs the unroller to unroll the loop `_value_` times. The parameter may optionally be enclosed in parentheses:

```
#pragma unroll 16
for (...) {
    ...
}

#pragma unroll(16)
for (...) {
    ...
}
```

Specifying `#pragma nounroll` indicates that the loop should not be unrolled:

```
#pragma nounroll
for (...) {
    ...
}
```

`#pragma unroll` and `#pragma unroll _value_` have identical semantics to `#pragma clang loop unroll(full)` and `#pragma clang loop unroll_count(_value_)` respectively. `#pragma nounroll` is equivalent to `#pragma clang loop unroll(disable)`. See [language extensions](#) for further details including limitations of the unroll hints.

### **`__attribute__((intel_reqd_sub_group_size))`**

#### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
X						X

The optional attribute `intel_reqd_sub_group_size` can be used to indicate that the kernel must be compiled and executed with the specified subgroup size. When this attribute is present, `get_max_sub_group_size()` is guaranteed to return the specified integer value. This is important for the correctness of many subgroup algorithms, and in some cases may be used by the compiler to generate more optimal code. See `cl_intel_required_subgroup_size` <[https://www.khronos.org/registry/OpenCL/extensions/intel/cl\\_intel\\_required\\_subgroup\\_size.txt](https://www.khronos.org/registry/OpenCL/extensions/intel/cl_intel_required_subgroup_size.txt)> for details.

### **`__attribute__((opencl_unroll_hint))`**

#### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
X						

The `opencl_unroll_hint` attribute qualifier can be used to specify that a loop (for, while and do loops) can be unrolled. This attribute qualifier can be used to specify full unrolling or partial unrolling by a specified amount. This is a compiler hint and the compiler may ignore this directive. See [OpenCL v2.0](#) s6.11.5 for details.

### **`__read_only, __write_only, __read_write (read_only, write_only, read_write)`**

#### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
				X		

The access qualifiers must be used with image object arguments or pipe arguments to declare if they are being read or written by a kernel or function.

The `read_only/_read_only`, `write_only/_write_only` and `read_write/_read_write` names are reserved for use as access qualifiers and shall not be used otherwise.

```
kernel void
foo (read_only image2d_t imageA,
     write_only image2d_t imageB) {
    ...
}
```

In the above example `imageA` is a read-only 2D image object, and `imageB` is a write-only 2D image object.

The `read_write` (or `_read_write`) qualifier can not be used with pipe.

More details can be found in the OpenCL C language Spec v2.0, Section 6.6.

## fallthrough, clang::fallthrough

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
	X	X				

The `fallthrough` (or `clang::fallthrough`) attribute is used to annotate intentional fall-through between switch labels. It can only be applied to a null statement placed at a point of execution between any statement and the next switch label. It is common to mark these places with a specific comment, but this attribute is meant to replace comments with a more strict annotation, which can be checked by the compiler. This attribute doesn't change semantics of the code and can be used wherever an intended fall-through occurs. It is designed to mimic control-flow statements like `break`;, so it can be placed in most places where `break`; can, but only if there are no statements on the execution path between it and the next switch label.

By default, Clang does not warn on unannotated `fallthrough` from one `switch` case to another. Diagnostics on `fallthrough` without a corresponding annotation can be enabled with the `-Wimplicit-fallthrough` argument.

Here is an example:

```
// compile with -Wimplicit-fallthrough
switch (n) {
case 22:
case 33: // no warning: no statements between case labels
    f();
case 44: // warning: unannotated fall-through
    g();
    [[clang::fallthrough]];
case 55: // no warning
    if (x) {
        h();
        break;
    }
    else {
        i();
        [[clang::fallthrough]];
    }
case 66: // no warning
    p();
    [[clang::fallthrough]]; // warning: fallthrough annotation does not
                           //      directly precede case label
    q();
case 77: // warning: unannotated fall-through
    r();
}
```

## suppress (gsl::suppress)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
	X					

The `[[gsl::suppress]]` attribute suppresses specific clang-tidy diagnostics for rules of the **C++ Core Guidelines** in a portable way. The attribute can be attached to declarations, statements, and at namespace scope.

```
[[gsl::suppress("Rh-public")]]
void f_() {
    int *p;
    [[gsl::suppress("type")]] {
        p = reinterpret_cast<int*>(7);
    }
}
namespace N {
    [[clang::suppress("type", "bounds")]];
    ...
}
```

## Calling Conventions

Clang supports several different calling conventions, depending on the target platform and architecture. The calling convention used for a function determines how parameters are passed, how results are returned to the caller, and other low-level details of calling a function.

### **fastcall** (`gnu::fastcall`, `__fastcall`, `_fastcall`)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X			X		

On 32-bit x86 targets, this attribute changes the calling convention of a function to use ECX and EDI as register parameters and clear parameters off of the stack on return. This convention does not support variadic calls or unprototyped functions in C, and has no effect on x86\_64 targets. This calling convention is supported primarily for compatibility with existing code. Users seeking register parameters should use the `regparm` attribute, which does not require callee-cleanup. See the documentation for **fastcall** on MSDN.

### **ms\_abi** (`gnu::ms_abi`)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

On non-Windows x86\_64 targets, this attribute changes the calling convention of a function to match the default convention used on Windows x86\_64. This attribute has no effect on Windows targets or non-x86\_64 targets.

**pcs (gnu::pcs)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

On ARM targets, this attribute can be used to select calling conventions similar to `stdcall` on x86. Valid parameter values are “aapcs” and “aapcs-vfp”.

**preserve\_all**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

On X86-64 and AArch64 targets, this attribute changes the calling convention of a function. The `preserve_all` calling convention attempts to make the code in the caller even less intrusive than the `preserve_most` calling convention. This calling convention also behaves identical to the `c` calling convention on how arguments and return values are passed, but it uses a different set of caller/callee-saved registers. This removes the burden of saving and recovering a large register set before and after the call in the caller. If the arguments are passed in callee-saved registers, then they will be preserved by the callee across the call. This doesn't apply for values returned in callee-saved registers.

On X86-64 the callee preserves all general purpose registers, except for R11. R11 can be used as a scratch register. Furthermore it also preserves all floating-point registers (XMMs/YMMs).



The idea behind this convention is to support calls to runtime functions that don't need to call out to any other functions.

This calling convention, like the `preserve_most` calling convention, will be used by a future version of the Objective-C runtime and should be considered experimental at this time.

## preserve\_most

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

On X86-64 and AArch64 targets, this attribute changes the calling convention of a function. The `preserve_most` calling convention attempts to make the code in the caller as unintrusive as possible. This convention behaves identically to the `c` calling convention on how arguments and return values are passed, but it uses a different set of caller/callee-saved registers. This alleviates the burden of saving and recovering a large register set before and after the call in the caller. If the arguments are passed in callee-saved registers, then they will be preserved by the callee across the call. This doesn't apply for values returned in callee-saved registers.

On X86-64 the callee preserves all general purpose registers, except for R11. R11 can be used as a scratch register. Floating-point registers (XMMs/YMMs) are not preserved and need to be saved by the caller.

The idea behind this convention is to support calls to runtime functions that have a hot path and a cold path. The hot path is usually a small piece of code that doesn't use many registers. The cold path might need to call out to another function and therefore only needs to preserve the caller-saved registers, which haven't already been saved by the caller. The *preserve\_most* calling convention is very similar to the *cold* calling convention in terms of caller/callee-saved registers, but they are used for different types of function calls. `coldcc` is for function calls that are rarely executed, whereas *preserve\_most* function calls are intended to be on the hot path and definitely executed a lot. Furthermore `preserve_most` doesn't prevent the inliner from inlining the function call.

This calling convention will be used by a future version of the Objective-C runtime and should therefore still be considered experimental at this time. Although this convention was created to optimize certain runtime calls to the Objective-C runtime, it is not limited to this runtime and might be used by other runtimes in the future too. The current implementation only supports X86-64 and AArch64, but the intention is to support more architectures in the future.

## regcall (gnu::regcall, \_\_regcall)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X			X		

On x86 targets, this attribute changes the calling convention to **\_\_regcall** convention. This convention aims to pass as many arguments as possible in registers. It also tries to utilize registers for the return value whenever it is possible.

**regparm (gnu::regparm)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

On 32-bit x86 targets, the regparm attribute causes the compiler to pass the first three integer parameters in EAX, EDX, and ECX instead of on the stack. This attribute has no effect on variadic functions, and all parameters are passed via the stack as normal.

**stdcall (gnu::stdcall, \_\_stdcall, \_stdcall)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X			X		

On 32-bit x86 targets, this attribute changes the calling convention of a function to clear parameters off of the stack on return. This convention does not support variadic calls or unprototyped functions in C, and has no effect on x86\_64 targets. This calling convention is used widely by the Windows API and COM applications. See the documentation for **\_\_stdcall** on MSDN.

**thiscall (gnu::thiscall, \_\_thiscall, \_thiscall)**

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X			X		

On 32-bit x86 targets, this attribute changes the calling convention of a function to use ECX for the first parameter (typically the implicit this parameter of C++ methods) and clear parameters off of the stack on return. This convention does not support variadic calls or unprototyped functions in C, and has no effect on x86\_64 targets. See the documentation for [\\_\\_thiscall](#) on MSDN.

## vectorcall (\_\_vectorcall, \_vectorcall)

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X				X		

On 32-bit x86 *and* x86\_64 targets, this attribute changes the calling convention of a function to pass vector parameters in SSE registers.

On 32-bit x86 targets, this calling convention is similar to `__fastcall`. The first two integer parameters are passed in ECX and EDX. Subsequent integer parameters are passed in memory, and callee clears the stack. On x86\_64 targets, the callee does *not* clear the stack, and integer parameters are passed in RCX, RDX, R8, and R9 as is done for the default Windows x64 calling convention.

On both 32-bit x86 and x86\_64 targets, vector and floating point arguments are passed in XMM0-XMM5. Homogeneous vector aggregates of up to four elements are passed in sequential SSE registers if enough are available. If AVX is enabled, 256 bit vectors are passed in YMM0-YMM5. Any vector or aggregate type that cannot be passed in registers for any reason is passed by reference, which allows the caller to align the parameter memory.

See the documentation for [\\_\\_vectorcall](#) on MSDN for more details.

## AMD GPU Attributes

### amdgpu\_flat\_work\_group\_size

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The flat work-group size is the number of work-items in the work-group size specified when the kernel is dispatched. It is the product of the sizes of the x, y, and z dimension of the work-group.

Clang supports the `__attribute__((amdgpu_flat_work_group_size(<min>, <max>)))` attribute for the AMDGPU target. This attribute may be attached to a kernel function definition and is an optimization hint.

<min> parameter specifies the minimum flat work-group size, and <max> parameter specifies the maximum flat work-group size (must be greater than <min>) to which all dispatches of the kernel will conform. Passing 0, 0 as <min>, <max> implies the default behavior (128, 256).

If specified, the AMDGPU target backend might be able to produce better machine code for barriers and perform scratch promotion by estimating available group segment size.

An error will be given if:

- Specified values violate subtarget specifications;
- Specified values are not compatible with values provided through other attributes.

## amdgpu\_num\_sgpr

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Clang supports the `__attribute__((amdgpu_num_sgpr(<num_sgpr>)))` and `__attribute__((amdgpu_num_vgpr(<num_vgpr>)))` attributes for the AMDGPU target. These attributes may be attached to a kernel function definition and are an optimization hint.

If these attributes are specified, then the AMDGPU target backend will attempt to limit the number of SGPRs and/or VGPRs used to the specified value(s). The number of used SGPRs and/or VGPRs may further be rounded up to satisfy the allocation requirements or constraints of the subtarget. Passing 0 as num\_sgpr and/or num\_vgpr implies the default behavior (no limits).

These attributes can be used to test the AMDGPU target backend. It is recommended that the `amdgpu_waves_per_eu` attribute be used to control resources such as SGPRs and VGPRs since it is aware of the limits for different subtargets.

An error will be given if:

- Specified values violate subtarget specifications;
- Specified values are not compatible with values provided through other attributes;
- The AMDGPU target backend is unable to create machine code that can meet the request.

## amdgpu\_num\_vgpr

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Clang supports the `__attribute__((amdgpu_num_sgpr(<num_sgpr>)))` and `__attribute__((amdgpu_num_vgpr(<num_vgpr>)))` attributes for the AMDGPU target. These attributes may be attached to a kernel function definition and are an optimization hint.

If these attributes are specified, then the AMDGPU target backend will attempt to limit the number of SGPRs and/or VGPRs used to the specified value(s). The number of used SGPRs and/or VGPRs may further be rounded up to satisfy the allocation requirements or constraints of the subtarget. Passing 0 as `num_sgpr` and/or `num_vgpr` implies the default behavior (no limits).

These attributes can be used to test the AMDGPU target backend. It is recommended that the `amdgpu_waves_per_eu` attribute be used to control resources such as SGPRs and VGPRs since it is aware of the limits for different subtargets.

An error will be given if:

- Specified values violate subtarget specifications;
- Specified values are not compatible with values provided through other attributes;
- The AMDGPU target backend is unable to create machine code that can meet the request.

## amdgpu\_waves\_per\_eu

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

A compute unit (CU) is responsible for executing the wavefronts of a work-group. It is composed of one or more execution units (EU), which are responsible for executing the wavefronts. An EU can have enough resources to maintain the state of more than one executing wavefront. This allows an EU to hide latency by switching between wavefronts in a similar way to symmetric multithreading on a CPU. In order to allow the state for multiple wavefronts to fit on an EU, the resources used by a single wavefront have to be limited. For example, the number of SGPRs and VGPRs. Limiting such resources can allow greater latency hiding, but can result in having to spill some register state to memory.

Clang supports the `__attribute__((amdgpu_waves_per_eu(<min>[, <max>])))` attribute for the AMDGPU target. This attribute may be attached to a kernel function definition and is an optimization hint.

`<min>` parameter specifies the requested minimum number of waves per EU, and *optional* `<max>` parameter specifies the requested maximum number of waves per EU (must be greater than `<min>` if specified). If `<max>` is omitted, then there is no restriction on the maximum number of waves per EU other than the one dictated by the hardware for which the kernel is compiled. Passing 0, 0 as `<min>`, `<max>` implies the default behavior (no limits).

If specified, this attribute allows an advanced developer to tune the number of wavefronts that are capable of fitting within the resources of an EU. The AMDGPU target backend can use this information to limit resources, such as number of SGPRs, number of VGPRs, size of available group and private memory segments, in such a way that guarantees that at least `<min>` wavefronts and at most `<max>` wavefronts are able to fit within the resources of an EU. Requesting more wavefronts can hide memory latency but limits available registers which can result in spilling. Requesting fewer wavefronts can help reduce cache thrashing, but can reduce memory latency hiding.

This attribute controls the machine code generated by the AMDGPU target backend to ensure it is capable of meeting the requested values. However, when the kernel is executed, there may be other reasons that prevent meeting the request, for example, there may be wavefronts from other kernels executing on the EU.

An error will be given if:

- Specified values violate subtarget specifications;
- Specified values are not compatible with values provided through other attributes;
- The AMDGPU target backend is unable to create machine code that can meet the request.

## Consumed Annotation Checking

Clang supports additional attributes for checking basic resource management properties, specifically for unique objects that have a single owning reference. The following attributes are currently supported, although **the implementation for these annotations is currently in development and are subject to change.**

## callable\_when

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Use `__attribute__((callable_when(...)))` to indicate what states a method may be called in. Valid states are unconsumed, consumed, or unknown. Each argument to this attribute must be a quoted string. E.g.:

```
__attribute__((callable_when("unconsumed", "unknown")))
```

## consumable

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Each class that uses any of the typestate annotations must first be marked using the `consumable` attribute. Failure to do so will result in a warning.

This attribute accepts a single parameter that must be one of the following: `unknown`, `consumed`, or `unconsumed`.

## param\_typestate

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

This attribute specifies expectations about function parameters. Calls to an function with annotated parameters will issue a warning if the corresponding argument isn't in the expected state. The attribute is also used to set the initial state of the parameter when analyzing the function's body.

## return\_tystate

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

The `return_tystate` attribute can be applied to functions or parameters. When applied to a function the attribute specifies the state of the returned value. The function's body is checked to ensure that it always returns a value in the specified state. On the caller side, values returned by the annotated function are initialized to the given state.

When applied to a function parameter it modifies the state of an argument after a call to the function returns. The function's body is checked to ensure that the parameter is in the expected state before returning.

## set\_tystate

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						X

Annotate methods that transition an object into a new state with `__attribute__((set_tystate(new_state)))`. The new state must be unconsumed, consumed, or unknown.

## test\_tystate

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute



X

X

Use `__attribute__((test_tystate(tested_state)))` to indicate that a method returns true if the object is in the specified state..

## Type Safety Checking

Clang supports additional attributes to enable checking type safety properties that can't be enforced by the C type system. To see warnings produced by these checks, ensure that `-Wtype-safety` is enabled. Use cases include:

- MPI library implementations, where these attributes enable checking that the buffer type matches the passed `MPI_Datatype`;
- for HDF5 library there is a similar use case to MPI;
- checking types of variadic functions' arguments for functions like `fcntl()` and `ioctl()`.

You can detect support for these attributes with `__has_attribute()`. For example:

```
#if defined(__has_attribute)
# if __has_attribute(argument_with_type_tag) && \
    __has_attribute(pointer_with_type_tag) && \
    __has_attribute(type_tag_for_datatype)
#   define ATTR_MPI_PWT(buffer_idx, type_idx) __attribute__((pointer_with_type_tag(mpi,buffer_idx,type_idx)))
/* ... other macros ... */
# endif
#endif

#if !defined(ATTR_MPI_PWT)
# define ATTR_MPI_PWT(buffer_idx, type_idx)
#endif

int MPI_Send(void *buf, int count, MPI_Datatype datatype /*, other args omitted */)
    ATTR_MPI_PWT(1,3);
```

## argument\_with\_type\_tag

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

Use `__attribute__((argument_with_type_tag(arg_kind, arg_idx, type_tag_idx)))` on a function declaration to specify that the function accepts a type tag that determines the type of some other argument.

This attribute is primarily useful for checking arguments of variadic functions (`pointer_with_type_tag` can be used in most non-variadic cases).

In the attribute prototype above:

`arg_kind` is an identifier that should be used when annotating all applicable type tags.

`arg_idx` provides the position of a function argument. The expected type of this function argument will be determined by the function argument specified by `type_tag_idx`. In the code example below, “3” means that the type of the function’s third argument will be determined by `type_tag_idx`.

`type_tag_idx` provides the position of a function argument. This function argument will be a type tag. The type tag will determine the expected type of the argument specified by `arg_idx`. In the code example below, “2” means that the type tag associated with the function’s second argument should agree with the type of the argument specified by `arg_idx`.

For example:

```
int fcntl(int fd, int cmd, ...)
    __attribute__((argument_with_type_tag(fcntl,3,2)));
// The function's second argument will be a type tag; this type tag will
// determine the expected type of the function's third argument.
```

## pointer\_with\_type\_tag

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

Use `__attribute__((pointer_with_type_tag(ptr_kind, ptr_idx, type_tag_idx)))` on a function declaration to specify that the function accepts a type tag that determines the pointee type of some other pointer argument.

In the attribute prototype above:

`ptr_kind` is an identifier that should be used when annotating all applicable type tags.

`ptr_idx` provides the position of a function argument; this function argument will have a pointer type. The expected pointee type of this pointer type will be determined by the function argument specified by `type_tag_idx`. In the code example below, “1” means that the pointee type of the function’s first argument will be determined by `type_tag_idx`.

`type_tag_idx` provides the position of a function argument; this function argument will be a type tag. The type tag will determine the expected pointee type of the pointer argument specified by `ptr_idx`. In the code example below, “3” means that the type tag associated with the function’s third argument should agree with the pointee type of the pointer argument specified by `ptr_idx`.

For example:

```
typedef int MPI_Datatype;
int MPI_Send(void *buf, int count, MPI_Datatype datatype /*, other args omitted */)
    __attribute__((pointer_with_type_tag(mpi,1,3)));
// The function's 3rd argument will be a type tag; this type tag will
// determine the expected pointee type of the function's 1st argument.
```

## type\_tag\_for\_datatype

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X						

When declaring a variable, use `__attribute__((type_tag_for_datatype(kind, type)))` to create a type tag that is tied to the type argument given to the attribute.

In the attribute prototype above:

`kind` is an identifier that should be used when annotating all applicable type tags.

`type` indicates the name of the type.

Clang supports annotating type tags of two forms.

**Type tag that is a reference to a declared identifier.** Use `__attribute__((type_tag_for_datatype(kind, type)))` when declaring that identifier:

```
typedef int MPI_Datatype;
extern struct mpi_datatype mpi_datatype_int
```

```
__attribute__(( type_tag_for_datatype(mpi,int) ));
#define MPI_INT ((MPI_Datatype) &mpi_datatype_int)
// &mpi_datatype_int is a type tag. It is tied to type "int".
```

**Type tag that is an integral literal.** Declare a static const variable with an initializer value and attach `__attribute__((type_tag_for_datatype(kind, type)))` on that declaration:

```
typedef int MPI_Datatype;
static const MPI_Datatype mpi_datatype_int
__attribute__(( type_tag_for_datatype(mpi,int) )) = 42;
#define MPI_INT ((MPI_Datatype) 42)
// The number 42 is a type tag. It is tied to type "int".
```

The `type_tag_for_datatype` attribute also accepts an optional third argument that determines how the type of the function argument specified by either `arg_idx` or `ptr_idx` is compared against the type associated with the type tag. (Recall that for the `argument_with_type_tag` attribute, the type of the function argument specified by `arg_idx` is compared against the type associated with the type tag. Also recall that for the `pointer_with_type_tag` attribute, the pointee type of the function argument specified by `ptr_idx` is compared against the type associated with the type tag.) There are two supported values for this optional third argument:

`layout_compatible` will cause types to be compared according to layout-compatibility rules (In C++11 [class.mem] p 17, 18, see the layout-compatibility rules for two standard-layout struct types and for two standard-layout union types). This is useful when creating a type tag associated with a struct or union type. For example:

```
/* In mpi.h */
typedef int MPI_Datatype;
struct internal_mpi_double_int { double d; int i; };
extern struct mpi_datatype mpi_datatype_double_int
__attribute__(( type_tag_for_datatype(mpi,
    struct internal_mpi_double_int, layout_compatible) ));

#define MPI_DOUBLE_INT ((MPI_Datatype) &mpi_datatype_double_int)

int MPI_Send(void *buf, int count, MPI_Datatype datatype, ...)
__attribute__(( pointer_with_type_tag(mpi,1,3) ));

/* In user code */
struct my_pair { double a; int b; };
struct my_pair *buffer;
MPI_Send(buffer, 1, MPI_DOUBLE_INT /*, ... */); // no warning because the
// layout of my_pair is
// compatible with that of
// internal_mpi_double_int
```

```

struct my_int_pair { int a; int b; }
struct my_int_pair *buffer2;
MPI_Send(buffer2, 1, MPI_DOUBLE_INT /*, ... */); // warning because the
// layout of my_int_pair
// does not match that of
// internal_mpi_double_int

```

`must_be_null` specifies that the function argument specified by either `arg_idx` (for the `argument_with_type_tag` attribute) or `ptr_idx` (for the `pointer_with_type_tag` attribute) should be a null pointer constant. The second argument to the `type_tag_for_datatype` attribute is ignored. For example:

```

/* In mpi.h */
typedef int MPI_Datatype;
extern struct mpi_datatype mpi_datatype_null
    __attribute__(( type_tag_for_datatype(mpi, void, must_be_null) ));

#define MPI_DATATYPE_NULL ((MPI_Datatype) &mpi_datatype_null)
int MPI_Send(void *buf, int count, MPI_Datatype datatype, ...)
    __attribute__(( pointer_with_type_tag(mpi, 1, 3) ));

/* In user code */
struct my_pair { double a; int b; };
struct my_pair *buffer;
MPI_Send(buffer, 1, MPI_DATATYPE_NULL /*, ... */); // warning: MPI_DATATYPE_NULL
// was specified but buffer
// is not a null pointer

```

## OpenCL Address Spaces

The address space qualifier may be used to specify the region of memory that is used to allocate the object. OpenCL supports the following address spaces: `__generic`(generic), `__global`(global), `__local`(local), `__private`(private), `__constant`(constant).

```

__constant int c = ...;

__generic int* foo(global int* g) {
    __local int* l;
    private int p;
    ...
}

```

```
return l;  
}
```

More details can be found in the OpenCL C language Spec v2.0, Section 6.5.

## constant (\_\_constant)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The constant address space attribute signals that an object is located in a constant (non-modifiable) memory region. It is available to all work items. Any type can be annotated with the constant address space attribute. Objects with the constant address space qualifier can be declared in any scope and must have an initializer.

## generic (\_\_generic)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The generic address space attribute is only available with OpenCL v2.0 and later. It can be used with pointer types. Variables in global and local scope and function parameters in non-kernel functions can have the generic address space type attribute. It is intended to be a placeholder for any other address space except for ‘\_\_constant’ in OpenCL code which can be used with multiple address spaces.

## global (\_\_global)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The global address space attribute specifies that an object is allocated in global memory, which is accessible by all work items. The content stored in this memory area persists between kernel executions. Pointer types to the global address space are allowed as function parameters or local variables. Starting with OpenCL v2.0, the global address space can be used with global (program scope) variables and static local variable as well.

## local (\_\_local)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The local address space specifies that an object is allocated in the local (work group) memory area, which is accessible to all work items in the same work group. The content stored in this memory region is not accessible after the kernel execution ends. In a kernel function scope, any variable can be in the local address space. In other scopes, only pointer types to the local address space are allowed. Local address space variables cannot have an initializer.

## private (\_\_private)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The private address space specifies that an object is allocated in the private (work item) memory. Other work items cannot access the same memory area and its content is destroyed after work item execution ends. Local variables can be declared in the private address space. Function arguments are always in the private address space. Kernel function arguments of a pointer or an array type cannot point to the private address space.

## Nullability Attributes

Whether a particular pointer may be “null” is an important concern when working with pointers in the C family of languages. The various nullability attributes indicate whether a particular pointer can be null or not, which makes APIs more expressive and can help static analysis

tools identify bugs involving null pointers. Clang supports several kinds of nullability attributes: the `nonnull` and `returns_nonnull` attributes indicate which function or method parameters and result types can never be null, while nullability type qualifiers indicate which pointer types can be null (`_Nullable`) or cannot be null (`_Nonnull`).

The nullability (type) qualifiers express whether a value of a given pointer type can be null (the `_Nullable` qualifier), doesn't have a defined meaning for null (the `_Nonnull` qualifier), or for which the purpose of null is unclear (the `_Null_unspecified` qualifier). Because nullability qualifiers are expressed within the type system, they are more general than the `nonnull` and `returns_nonnull` attributes, allowing one to express (for example) a nullable pointer to an array of nonnull pointers. Nullability qualifiers are written to the right of the pointer to which they apply. For example:

```
// No meaningful result when 'ptr' is null (here, it happens to be undefined behavior).
int fetch(int * _Nonnull ptr) { return *ptr; }

// 'ptr' may be null.
int fetch_or_zero(int * _Nullable ptr) {
    return ptr ? *ptr : 0;
}

// A nullable pointer to non-null pointers to const characters.
const char *join_strings(const char * _Nonnull * _Nullable strings, unsigned n);
```

In Objective-C, there is an alternate spelling for the nullability qualifiers that can be used in Objective-C methods and properties using context-sensitive, non-underscored keywords. For example:

```
@interface NSView : NSResponder
- (nullable NSView *)ancestorSharedWithView:(nonnull NSView *)aView;
@property (assign, nullable) NSView *superview;
@property (readonly, nonnull) NSArray *subviews;
@end
```

## **`_Nonnull`**

### Supported Syntaxes

GNU	C++11	C2x	<code>__declspec</code>	Keyword	Pragma	Pragma clang attribute
-----	-------	-----	-------------------------	---------	--------	------------------------



X

The `_Nonnull` nullability qualifier indicates that null is not a meaningful value for a value of the `_Nonnull` pointer type. For example, given a declaration such as:

```
int fetch(int * _Nonnull ptr);
```

a caller of `fetch` should not provide a null value, and the compiler will produce a warning if it sees a literal null value passed to `fetch`. Note that, unlike the declaration attribute `nonnull`, the presence of `_Nonnull` does not imply that passing null is undefined behavior: `fetch` is free to consider null undefined behavior or (perhaps for backward-compatibility reasons) defensively handle null.

## **\_Null\_unspecified**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The `_Null_unspecified` nullability qualifier indicates that neither the `_Nonnull` nor `_Nullable` qualifiers make sense for a particular pointer type. It is used primarily to indicate that the role of null with specific pointers in a nullability-annotated header is unclear, e.g., due to overly-complex implementations or historical factors with a long-lived API.

## **\_Nullable**

Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
				X		

The `_Nullable` nullability qualifier indicates that a value of the `_Nullable` pointer type can be null. For example, given:

```
int fetch_or_zero(int * _Nullable ptr);
```

a caller of `fetch_or_zero` can provide null.

## nonnull (gnu::nonnull)

### Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					

The `nonnull` attribute indicates that some function parameters must not be null, and can be used in several different ways. It's original usage (**from GCC**) is as a function (or Objective-C method) attribute that specifies which parameters of the function are nonnull in a comma-separated list. For example:

```
extern void * my_memcpy (void *dest, const void *src, size_t len)
    __attribute__((nonnull (1, 2)));
```

Here, the `nonnull` attribute indicates that parameters 1 and 2 cannot have a null value. Omitting the parenthesized list of parameter indices means that all parameters of pointer type cannot be null:

```
extern void * my_memcpy (void *dest, const void *src, size_t len)
    __attribute__((nonnull));
```

Clang also allows the `nonnull` attribute to be placed directly on a function (or Objective-C method) parameter, eliminating the need to specify the parameter index ahead of type. For example:

```
extern void * my_memcpy (void *dest __attribute__((nonnull)),
    const void *src __attribute__((nonnull)), size_t len);
```

Note that the `nonnull` attribute indicates that passing null to a non-null parameter is undefined behavior, which the optimizer may take advantage of to, e.g., remove null checks. The `_Nonnull` type qualifier indicates that a pointer cannot be null in a more general manner (because it is part of the type system) and does not imply undefined behavior, making it more widely applicable.

## returns\_nonnull (gnu::returns\_nonnull)

## Supported Syntaxes

GNU	C++11	C2x	__declspec	Keyword	Pragma	Pragma clang attribute
X	X					X

The `returns_nonnull` attribute indicates that a particular function (or Objective-C method) always returns a non-null pointer. For example, a particular system `malloc` might be defined to terminate a process when memory is not available rather than returning a null pointer:

```
extern void * malloc (size_t size) __attribute__((returns_nonnull));
```

The `returns_nonnull` attribute implies that returning a null pointer is undefined behavior, which the optimizer may take advantage of. The `_Nonnull` type qualifier indicates that a pointer cannot be null in a more general manner (because it is part of the type system) and does not imply undefined behavior, making it more widely applicable