

Performance Testing

Android O includes binder and hwbinder performance tests for throughput and latency. While many scenarios exist for detecting perceptible performance problems, running such scenarios can be time consuming and results are often unavailable until after a system is integrated. Using the provided performance tests makes it easier to test during development, detect serious problems earlier, and improve user experience.

Performance tests include the following four categories:

- binder throughput (available in `system/libhwbinder/vts/performance/Benchmark_binder.cpp`)
- binder latency (available in `frameworks/native/libs/binder/tests/schd-dbg.cpp`)
- hwbinder throughput (available in `system/libhwbinder/vts/performance/Benchmark.cpp`)
- hwbinder latency (available in `system/libhwbinder/vts/performance/Latency.cpp`)

About binder and hwbinder

Binder and hwbinder are Android inter-process communication (IPC) infrastructures that share the same Linux driver but have the following qualitative differences:

Aspect	binder	hwbinder
Purpose	Provide a general purpose IPC scheme for framework	Communicate with hardware
Property	Optimized for Android framework usage	Minimum overhead low latency
Change scheduling policy for foreground/background	Yes	No
Arguments passing	Uses serialization supported by Parcel object	Uses scatter buffers and avoids the overhead to copy data required for Parcel serialization
Priority inheritance	No	Yes

Binder and hwbinder processes

A systrace visualizer displays transactions as follows:

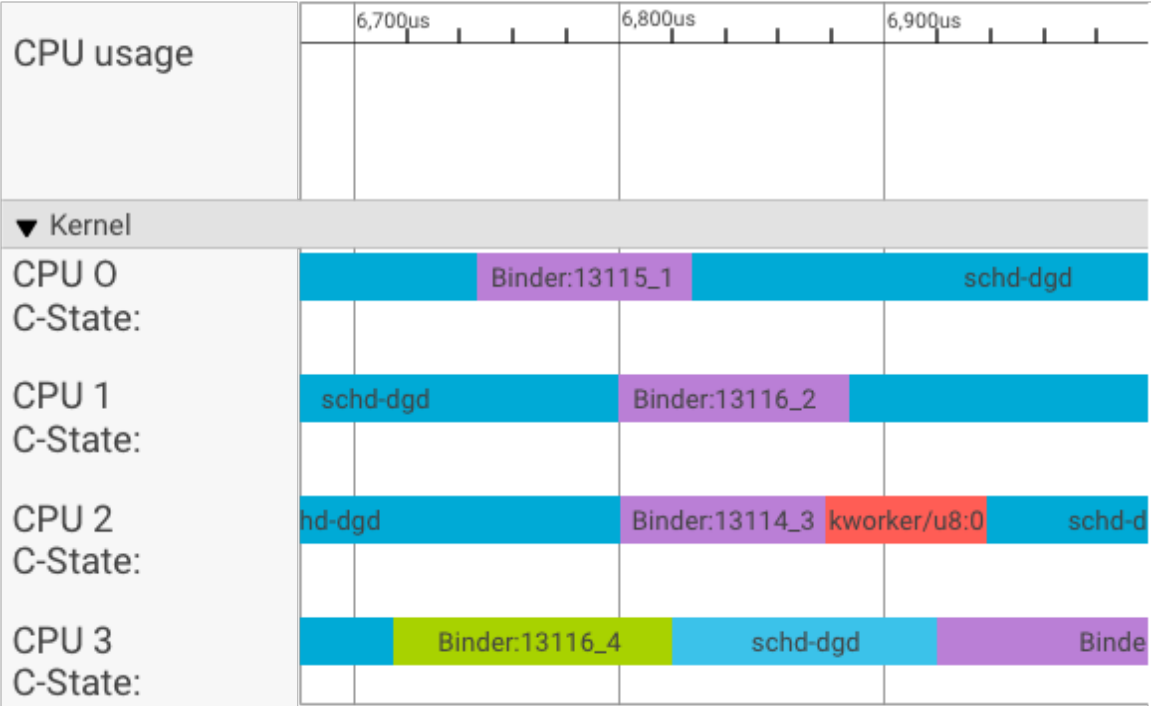


Figure 1. Systrace visualization of binder processes.

In the above example:

- The four (4) schd-dbg processes are client processes.
- The four (4) binder processes are server processes (name starts with **Binder** and ends with a sequence number).

- A client process is always paired with a server process, which is dedicated to its client.
- All the client-server process pairs are scheduled independently by kernel concurrently.

In CPU 1, the OS kernel executes the client to issue the request. It then uses the same CPU whenever possible to wake up a server process, handle the request, and context switch back after the request is complete.

Throughput vs. latency

In a perfect transaction, where the client and server process switch seamlessly, throughput and latency tests do not produce substantially different messages. However, when the OS kernel is handling an interrupt request (IRQ) from hardware, waiting for locks, or simply choosing not to handle a message immediately, a latency bubble can form.

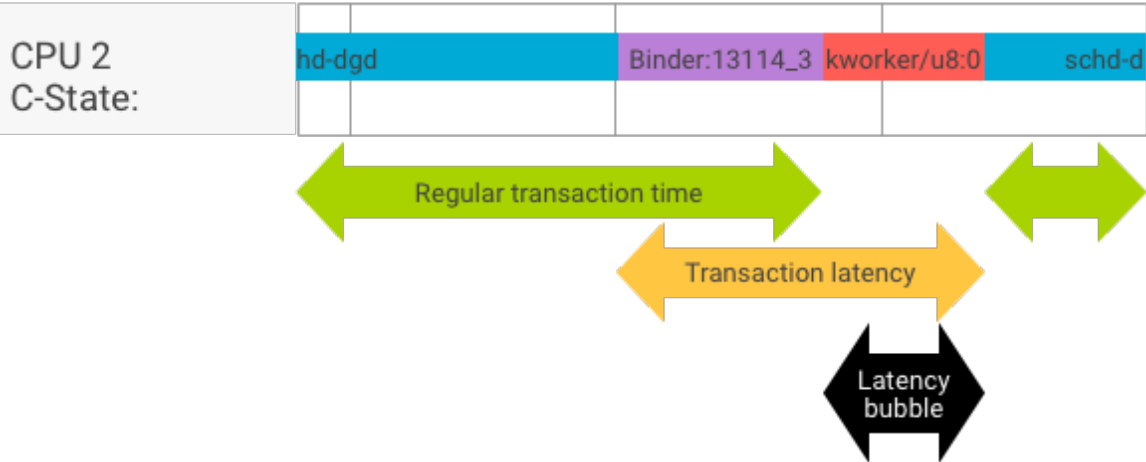


Figure 2. Latency bubble due to differences in throughput and latency.

The throughput test generates a large number of transactions with different payload sizes, providing a good estimation for the regular transaction time (in best case scenarios) and the maximum throughput the binder can achieve.

In contrast, the latency test performs no actions on the payload to minimize the regular transaction time. We can use transaction time to estimate the binder overhead, make statistics for the worst case, and calculate the ratio of transactions whose latency meets a specified deadline.

Handling priority inversions

A priority inversion occurs when a thread with higher priority is logically waiting for a thread with lower priority. Real-time (RT) applications have a priority inversion problem:




Figure 3. Priority inversion in real-time applications.

When using Linux Completely Fair Scheduler (CFS) scheduling, a thread always has a chance to run even when other threads have a higher priority. As a result, applications with CFS scheduling handle priority inversion as expected behavior and not as a problem. In cases where the Android framework needs RT scheduling to guarantee the privilege of high priority threads however, priority inversion must be resolved.

Example priority inversion during a binder transaction (RT thread is logically blocked by other CFS threads when waiting for a binder thread to service):

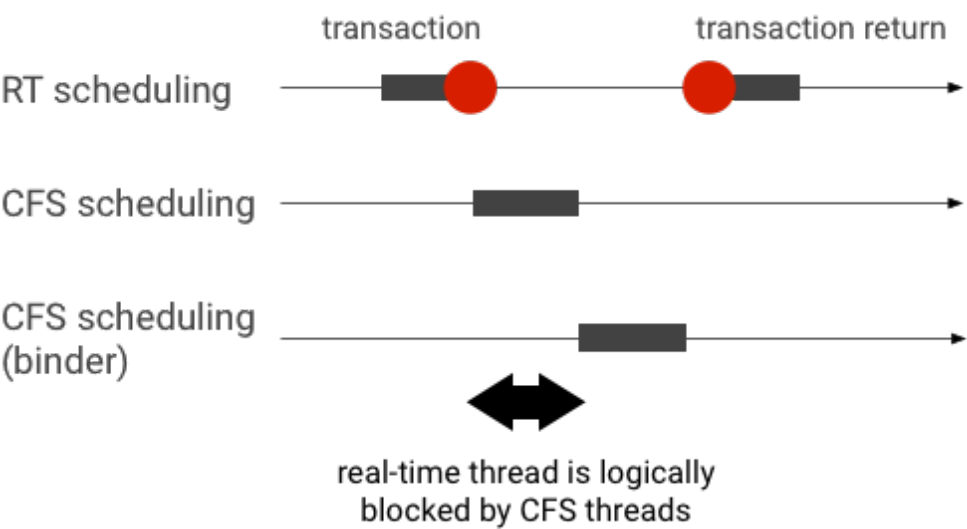


Figure 4. Priority inversion, blocked real-time threads.

To avoid blockages, you can use priority inheritance to temporarily escalate the Binder thread to a RT thread when it services a request from a RT client. Keep in mind that RT scheduling has limited resources and should be used carefully. In a system with n CPUs, the maximum number of current RT threads is also n ; additional RT threads might need to wait (and thus miss their deadlines) if all CPUs are taken by other RT threads.

To resolve all possible priority inversions, you could use priority inheritance for both binder and hwbinder. However, as binder is widely used across the system, enabling priority inheritance for binder transactions might spam the system with more RT threads than it can service.

Running throughput tests

The throughput test is run against binder/hwbinder transaction throughput. In a system that is not overloaded, latency bubbles are rare and their impact can be eliminated as long as the number of iterations is high enough.

- The **binder** throughput test is in `system/libhwbinder/vts/performance/Benchmark_binder.cpp`.
- The **hwbinder** throughput test is in `system/libhwbinder/vts/performance/Benchmark.cpp`.

Test results

Example throughput test results for transactions using different payload sizes:

Benchmark	Time	CPU	Iterations
BM_sendVec_binderize/4	70302 ns	32820 ns	21054
BM_sendVec_binderize/8	69974 ns	32700 ns	21296
BM_sendVec_binderize/16	70079 ns	32750 ns	21365
BM_sendVec_binderize/32	69907 ns	32686 ns	21310
BM_sendVec_binderize/64	70338 ns	32810 ns	21398
BM_sendVec_binderize/128	70012 ns	32768 ns	21377
BM_sendVec_binderize/256	69836 ns	32740 ns	21329
BM_sendVec_binderize/512	69986 ns	32830 ns	21296
BM_sendVec_binderize/1024	69714 ns	32757 ns	21319
BM_sendVec_binderize/2k	75002 ns	34520 ns	20305
BM_sendVec_binderize/4k	81955 ns	39116 ns	17895
BM_sendVec_binderize/8k	95316 ns	45710 ns	15350
BM_sendVec_binderize/16k	112751 ns	54417 ns	12679
BM_sendVec_binderize/32k	146642 ns	71339 ns	9901
BM_sendVec_binderize/64k	214796 ns	104665 ns	6495

- **Time** indicates the round trip delay measured in real time.
- **CPU** indicates the accumulated time when CPUs are scheduled for the test.
- **Iterations** indicates the number of times the test function executed.

For example, for an 8-byte payload:

BM_sendVec_binderize/8	69974 ns	32700 ns	21296
------------------------	----------	----------	-------

... the maximum throughput the binder can achieve is calculated as:

MAX throughput with 8-byte payload = $(8 * 21296) / 69974 \approx 2.423 \text{ b/ns} \approx 2.268 \text{ Gb/s}$

Test options

To get results in .json, run the test with the `--benchmark_format=json` argument:

```
$ libhwbinder_benchmark --benchmark_format=json
{
  "context": {
    "date": "2017-05-17 08:32:47",
    "num_cpus": 4,
    "mhz_per_cpu": 19,
    "cpu_scaling_enabled": true,
    "library_build_type": "release"
  },
  "benchmarks": [
    {
      "name": "BM_sendVec_binderize/4",
      "iterations": 32342,
      "real_time": 47809,
      "cpu_time": 21906,
      "time_unit": "ns"
    },
    ...
  ]
}
```

Running latency tests

The latency test measures the time it takes for the client to begin initializing the transaction, switch to the server process for handling, and receive the result. The test also looks for known bad scheduler behaviors that can negatively impact transaction latency, such as a scheduler that does not support priority inheritance or honor the sync flag.

- The binder latency test is in `frameworks/native/libs/binder/tests/schd-dbg.cpp`.
- The hwbinder latency test is in `system/libhwbinder/vts/performance/Latency.cpp`.

Test results

Results (in .json) show statistics for average/best/worst latency and the number of deadlines missed.

Test options

Latency tests take the following options:

Command	Description
<code>-i value</code>	Specify number of iterations.
<code>-pair value</code>	Specify the number of process pairs.
<code>-deadline_us 2500</code>	Specify the deadline in us.
<code>-v</code>	Get verbose (debugging) output.
<code>-trace</code>	Halt the trace on a deadline hit.

The following sections detail each option, describe usage, and provide example results.

Specifying iterations

Example with a large number of iterations and verbose output disabled:

```
$ libhwbinder_latency -i 5000 -pair 3
{
```

```
"cfg":{"pair":3,"iterations":5000,"deadline_us":2500},
"P0":{"SYNC":"GOOD","S":9352,"I":10000,"R":0.9352,
  "other_ms":{"avg":0.2 , "wst":2.8 , "bst":0.053, "miss":2, "meetR":0.9996},
  "fifo_ms": { "avg":0.16, "wst":1.5 , "bst":0.067, "miss":0, "meetR":1}
},
"P1":{"SYNC":"GOOD","S":9334,"I":10000,"R":0.9334,
  "other_ms":{"avg":0.19, "wst":2.9 , "bst":0.055, "miss":2, "meetR":0.9996},
  "fifo_ms": { "avg":0.16, "wst":3.1 , "bst":0.066, "miss":1, "meetR":0.9998}
},
"P2":{"SYNC":"GOOD","S":9369,"I":10000,"R":0.9369,
  "other_ms":{"avg":0.19, "wst":4.8 , "bst":0.055, "miss":6, "meetR":0.9988},
  "fifo_ms": { "avg":0.15, "wst":1.8 , "bst":0.067, "miss":0, "meetR":1}
},
"inheritance": "PASS"
}
```

These test results show the following:

"pair":3

Creates one client and server pair.

"iterations": 5000

Includes 5000 iterations.

"deadline_us":2500

Deadline is 2500us (2.5ms); most transactions are expected to meet this value.

"I": 10000

A single test iteration includes two (2) transactions:

- One transaction by normal priority (CFS `other`)
- One transaction by real time priority (RT-`fifo`)

5000 iterations equals a total of 10000 transactions.

"S": 9352

9352 of the transactions are synced in the same CPU.

"R": 0.9352

Indicates the ratio at which the client and server are synced together in the same CPU.

"other_ms":{"avg":0.2 , "wst":2.8 , "bst":0.053, "miss":2, "meetR":0.9996}

The average (`avg`), worst (`wst`), and the best (`bst`) case for all transactions issued by a normal priority caller. Two transactions miss the deadline, making the meet ratio (`meetR`) 0.9996.

"fifo_ms": { "avg":0.16, "wst":1.5 , "bst":0.067, "miss":0, "meetR":1}

Similar to `other_ms`, but for transactions issued by client with `rt_fifo` priority. It's likely (but not required) that the `fifo_ms` has a better result than `other_ms`, with lower `avg` and `wst` values and a higher `meetR` (the difference can be even more significant with load in the background).

Note: Background load may impact the throughput result and the `other_ms` tuple in the latency test. Only the `fifo_ms` may show similar results as long as the background load has a lower priority than `RT-fifo`.

Specifying pair values

Each client process is paired with a server process dedicated for the client, and each pair may be scheduled independently to any CPU. However, the CPU migration should not happen during a transaction as long as the SYNC flag is `honor`.

Ensure the system is not overloaded! While high latency in an overloaded system is expected, test results for an overloaded system do not provide useful information. To test a system with higher pressure, use `-pair #cpu-1` (or `-pair #cpu` with caution). Testing using `-pair n`

with $n > \text{\#cpu}$ overloads the system and generates useless information.

Specifying deadline values

After extensive user scenario testing (running the latency test on a qualified product), we determined that 2.5ms is the deadline to meet. For new applications with higher requirements (such as 1000 photos/second), this deadline value will change.

Specifying verbose output

Using the `-v` option displays verbose output. Example:

```
$ libhwbinder_latency -i 1 -v

-----
service      pid: 8674 tid: 8674 cpu: 1
SCHED_OTHER 0

-----

main         pid: 8673 tid: 8673 cpu: 1

-----

client       pid: 8677 tid: 8677 cpu: 0
SCHED_OTHER 0

-----

fifo-caller  pid:8677 tid: 8678 cpu: 0
SCHED_FIFO  99

-----

hwbinder     pid: 8674 tid: 8676 cpu: 0
???         99

-----

other-caller pid: 8677 tid: 8677 cpu: 0
SCHED_OTHER 0

-----

hwbinder     pid: 8674 tid: 8676 cpu: 0
SCHED_OTHER 0
```

- The **service thread** is created with a `SCHED_OTHER` priority and run in `CPU:1` with `pid 8674`.
- The **first transaction** is then started by a `fifo-caller`. To service this transaction, the hwbinder upgrades the priority of server (`pid: 8674 tid: 8676`) to be 99 and also marks it with a transient scheduling class (printed as `???`). The scheduler then puts the server process in `CPU:0` to run and syncs it with the same CPU with its client.
- The **second transaction** caller has a `SCHED_OTHER` priority. The server downgrades itself and services the caller with `SCHED_OTHER` priority.

Using trace for debugging

You can specify the `-t trace` option to debug latency issues. When used, the latency test stops the tracelog recording at the moment when bad latency is detected. Example:

```
$ atrace --async_start -b 8000 -c sched idle workq binder_driver sync freq
$ libhwbinder_latency -deadline_us 50000 -trace -i 50000 -pair 3
deadline triggered: halt ± stop trace
log:/sys/kernel/debug/tracing/trace
```

The following components can impact latency:

- **Android build mode.** Eng mode is usually slower than userdebug mode.
- **Framework.** How does the framework service use `ioctl` to config to the binder?
- **Binder driver.** Does the driver support fine-grained locking? Does the driver contain all performance turning patches?

- **Kernel version.** The better real time capability the kernel has, the better the results.
- **Kernel config.** Does the kernel config contain DEBUG configs such as DEBUG_PREEMPT and DEBUG_SPIN_LOCK?
- **Kernel scheduler.** Does the kernel have an Energy-Aware scheduler (EAS) or Heterogeneous Multi-Processing (HMP) scheduler? Are there kernel drivers (cpu-freq driver, cpu-idle driver, cpu-hotplug, etc.) that impact the scheduler?

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](http://creativecommons.org/licenses/by/3.0/) (<http://creativecommons.org/licenses/by/3.0/>), and code samples are licensed under the [Apache 2.0 License](http://www.apache.org/licenses/LICENSE-2.0) (<http://www.apache.org/licenses/LICENSE-2.0>). For details, see our [Site Policies](https://developers.google.com/terms/site-policies) (<https://developers.google.com/terms/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated August 21, 2017.