




Android Development

≡ Primary Menu

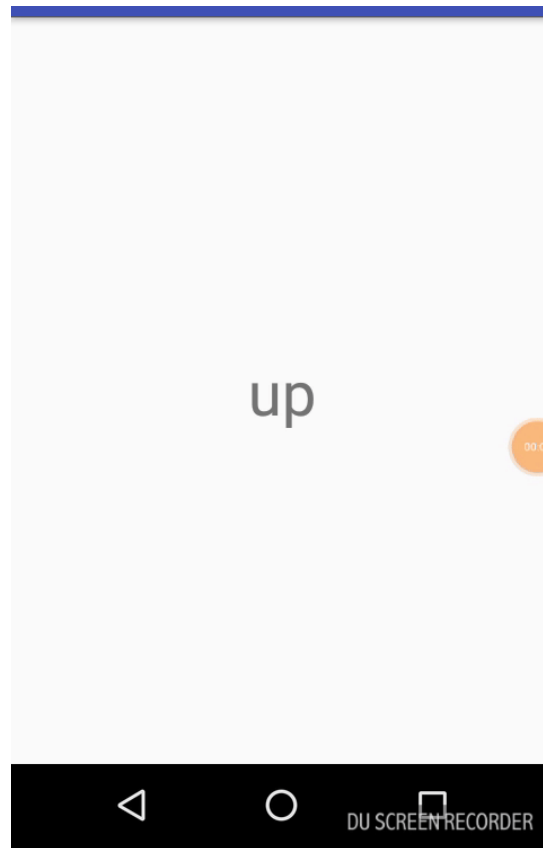
Speech Recognition Using TensorFlow

 brijeshthumar /  September 10, 2017 /  TensorFlow

This tutorial will show you how to runs a simple speech recognition model built by the [audio training tutorial](#). Listens for a small set of words, and display them in the UI when they are recognized.

It's important to know that real speech and audio recognition systems are much more complex, but like [MNIST for images](#), it should give you a basic understanding of the techniques involved. Once you've completed this tutorial, you'll have a application that tries to classify a one second audio clip as either silence, an unknown word, "yes", "no", "up", "down", "left", "right", "on", "off", "stop", or "go".





1.Preparation

You can train your model on the desktop or on the laptop or on the server and then you can use that pre-trained model on our mobile device. So there's no training that would happen on the device the training would happen on our bigger machine either a server or our laptop. You can download [a pretrained model from tensorflow.org](https://www.tensorflow.org/)

2. Adding Dependencies

The TensorFlow Inference Interface is available as a [JCenter package](#) and can be included quite simply in your android project with a couple of lines

in the project's build.gradle file:

```
1 allprojects {  
2     repositories {  
3         jcenter()  
4     }  
5 }
```

Add the following dependency in app's build.gradle

```
1 dependencies {  
2     .....  
3     compile 'org.tensorflow:tensorflow-android:+'  
4 }
```

This will tell Gradle to use the latest version of the TensorFlow AAR that has been released to <https://bintray.com/google/tensorflow/tensorflow-android>. You may replace the + with an explicit version label if you wish to use a specific release of TensorFlow in your app.

3.Add Pre-trained Model to Project

You need the **pre-trained** model and **label** file.You can download the model from here.Unzip this zip file, You will get `conv_actions_labels.txt` (label for objects) and `conv_actions_frozen.pb` (pre-trained model).

Put `conv_actions_labels.txt` and `conv_actions_frozen.pb` into `android/assets` directory.

4.Microphone Permission

To request microphone, you should be requesting `RECORD_AUDIO`

permission in your manifest file as below:

```
1 <uses-permission android:name="android.permission.RECORD_AUD
```

Since Android 6.0 Marshmallow, the application will not be granted any permission at installation time. Instead, the application has to ask the user for a permission one-by-one at runtime.

```
1 private void requestMicrophonePermission() {  
2     ActivityCompat.requestPermissions(MainActivity.this  
3         new String[]{android.Manifest.permission.RE  
4     }  
5     @Override  
6     public void onRequestPermissionsResult(int requestCode, Str  
7         if (requestCode == REQUEST_RECORD_AUDIO&& grantResults  
8             && grantResults[0] == PackageManager.PERMIS  
9             startRecording();  
10            startRecognition();  
11        }  
12    }
```

5.Recording Audio

The AudioRecord class manages the audio resources for Java applications to record audio from the audio input hardware of the platform. This is achieved by “pulling” (reading) the data from the AudioRecord object. The application is responsible for polling the AudioRecord object in time using read(short[], int, int).

```
1 private void record() {  
2     android.os.Process.setThreadPriority(android.os.Pro  
3  
4     // Estimate the buffer size we'll need for this dev  
5     int bufferSize =  
6         AudioRecord.getMinBufferSize(  
7             SAMPLE_RATE, AudioFormat.CHANNEL_IN
```

```
8      if (bufferSize == AudioRecord.ERROR || bufferSize =
9          bufferSize = SAMPLE_RATE * 2;
10     }
11     short[] audioBuffer = new short[bufferSize / 2];
12
13     AudioRecord record =
14         new AudioRecord(
15             MediaRecorder.AudioSource.DEFAULT,
16             SAMPLE_RATE,
17             AudioFormat.CHANNEL_IN_MONO,
18             AudioFormat.ENCODING_PCM_16BIT,
19             bufferSize);
20
21     if (record.getState() != AudioRecord.STATE_INITIALI
22         Log.e(LOG_TAG, "Audio Record can't initialize!"
23         return;
24     }
25
26     record.startRecording();
27
28     Log.v(LOG_TAG, "Start recording");
29
30     // Loop, gathering audio data and copying it to a r
31     while (shouldContinue) {
32         int numberRead = record.read(audioBuffer, 0, au
33         int maxLength = recordingBuffer.length;
34         int newRecordingOffset = recordingOffset + numb
35         int secondCopyLength = Math.max(0, newRecording
36         int firstCopyLength = numberRead - secondCopyLe
37         // We store off all the data for the recognitio
38         // thread will copy out of this buffer into its
39         // lock, so this should be thread safe.
40         recordingBufferLock.lock();
41         try {
42             System.arraycopy(audioBuffer, 0, recordingB
43             System.arraycopy(audioBuffer, firstCopyLeng
44             recordingOffset = newRecordingOffset % maxL
45         } finally {
46             recordingBufferLock.unlock();
47         }
48     }
49
50     record.stop();
51     record.release();
52 }
```

6.Run TensorFlow Model

A `TensorFlowInferenceInterface` class that provides a smaller API surface suitable for inference and summarizing the performance of model execution.

```
1 private void recognize() {
2     Log.v(LOG_TAG, "Start recognition");
3
4     short[] inputBuffer = new short[RECORDING_LENGTH];
5     float[] floatInputBuffer = new float[RECORDING_LENGTH];
6     float[] outputScores = new float[labels.size()];
7     String[] outputScoresNames = new String[]{OUTPUT_SCORES
8     int[] sampleRateList = new int[]{SAMPLE_RATE};
9
10    // Loop, grabbing recorded data and running the recogni
11    while (shouldContinueRecognition) {
12        // The recording thread places data in this rou
13        // make sure there's no writing happening and t
14        // local version.
15        recordingBufferLock.lock();
16        try {
17            int maxLength = recordingBuffer.length;
18            int firstCopyLength = maxLength - recordingOffse
19            int secondCopyLength = recordingOffset;
20            System.arraycopy(recordingBuffer, recordingOffse
21            System.arraycopy(recordingBuffer, 0, inputBuffer
22        } finally {
23            recordingBufferLock.unlock();
24        }
25
26        // We need to feed in float values between -1.0
27        // signed 16-bit inputs.
28        for (int i = 0; i < RECORDING_LENGTH; ++i) {
29            floatInputBuffer[i] = inputBuffer[i] / 32767.0
30        }
31
32        // Run the model.
33        inferenceInterface.feed(SAMPLE_RATE_NAME, sampleRate
34        inferenceInterface.feed(INPUT_DATA_NAME, floatInputB
35        inferenceInterface.run(outputScoresNames);
36        inferenceInterface.fetch(OUTPUT_SCORES_NAME, outputsS
37    }
```

```

38      // Use the smoother to figure out if we've had a rea
39      long currentTime = System.currentTimeMillis();
40      final RecognizeCommands.RecognitionResult result = r
41
42      runOnUiThread(
43          new Runnable() {
44              @Override
45              public void run() {
46                  // If we do have a new command, highligh
47                  if (!result.foundCommand.startsWith("
48                      int labelIndex = -1;
49                      for (int i = 0; i < labels.size
50                      if (labels.get(i).equals(resul
51                          labelIndex = i;
52                      }
53                  }
54                  label.setText(result.foundCommand)
55              }
56          });
57      try {
58          // We don't need to run too frequently, so snoo
59          Thread.sleep(MINIMUM_TIME_BETWEEN_SAMPLES_MS);
60      } catch (InterruptedException e) {
61          // Ignore
62      }
63  }
64  }
65
66  Log.v(LOG_TAG, "End recognition");
67 }

```

7. Recognize Commands

RecognizeCommands class is fed the output of running the TensorFlow model over time, it averages the signals and returns information about a label when it has enough evidence to think that a recognized word has been found. The implementation is fairly small, just keeping track of the last few predictions and averaging them.

```

1 public RecognitionResult processLatestResults(float[] curre
2     if (currentResults.length != labelsCount) {

```

```
3         throw new RuntimeException(  
4             "The results for recognition should con  
5                 + labelsCount  
6                 + " elements, but there are "  
7                 + currentResults.length);  
8     }  
9  
10    if ((!previousResults.isEmpty()) && (currentTimeMS  
11        throw new RuntimeException(  
12            "You must feed results in increasing ti  
13                + currentTimeMS  
14                + " that was earlier than the p  
15                + previousResults.getFirst().fi  
16    }  
17  
18    final int howManyResults = previousResults.size();  
19    // Ignore any results that are coming in too freque  
20    if (howManyResults > 1) {  
21        final long timeSinceMostRecent = currentTimeMS  
22        if (timeSinceMostRecent < minimumTimeBetweenSam  
23            return new RecognitionResult(previousTopLab  
24    }  
25    }  
26  
27    // Add the latest results to the head of the queue.  
28    previousResults.addLast(new Pair<Long, float[]>(cur  
29    Log.d(TAG, currentResults + " " + currentTimeMS);  
30  
31    // Prune any earlier results that are too old for t  
32    final long timeLimit = currentTimeMS - averageWindo  
33    while (previousResults.getFirst().first < timeLimit  
34        previousResults.removeFirst();  
35    }  
36    // If there are too few results, assume the result  
37    // bail.  
38    final long earliestTime = previousResults.getFirst(  
39    final long samplesDuration = currentTimeMS - earlie  
40    if ((howManyResults < minimumCount)  
41        || (samplesDuration < (averageWindowDuratio  
42        Log.v("RecognizeResult", "Too few results");  
43        return new RecognitionResult(previousTopLabel,  
44    }  
45  
46    // Calculate the average score across all the resul  
47    float[] averageScores = new float[labelsCount];  
48    for (Pair<Long, float[]> previousResult : previousR  
49        final float[] scoresTensor = previousResult.sec  
50        int i = 0;  
51        while (i < scoresTensor.length) {
```



```
52         averageScores[i] += scoresTensor[i] / howMa
53         ++i;
54     }
55 }
56
57 // Sort the averaged results in descending score or
58 ScoreForSorting[] sortedAverageScores = new ScoreFo
59 for (int i = 0; i < labelsCount; ++i) {
60     sortedAverageScores[i] = new ScoreForSorting(av
61 }
62 Arrays.sort(sortedAverageScores);
63
64 // See if the latest top score is enough to trigger
65 final int currentTopIndex = sortedAverageScores[0].
66 final String currentTopLabel = labels.get(currentTo
67 final float currentTopScore = sortedAverageScores[0
68 // If we've recently had another label trigger, ass
69 // soon afterwards is a bad result.
70 long timeSinceLastTop;
71 if (previousTopLabel.equals(SILENCE_LABEL) || (prev
72     timeSinceLastTop = Long.MAX_VALUE;
73 } else {
74     timeSinceLastTop = currentTimeMS - previousTopL
75 }
76 boolean isNewCommand;
77 if ((currentTopScore > detectionThreshold) && (time
78     previousTopLabel = currentTopLabel;
79     previousTopLabelTime = currentTimeMS;
80     previousTopLabelScore = currentTopScore;
81     isNewCommand = true;
82 } else {
83     isNewCommand = false;
84 }
85 return new RecognitionResult(currentTopLabel, curre
86 }
```

The demo app updates its UI of results automatically based on the labels text file you copy into assets alongside your frozen graph, which means you can easily try out different models without needing to make any code changes. You will need to update `LABEL_FILENAME` and `MODEL_FILENAME` to point to the files you've added if you change the paths though.

8.conclusion

You can easily replace it with a model you've trained yourself. If you do this, you'll need to make sure that the constants in [the main MainActivity Java source file](#) like `SAMPLE_RATE` and `SAMPLE_DURATION` match any changes you've made to the defaults while training. You'll also see that there's a [Java version of the RecognizeCommands module](#) that's very similar to the C++ version in this tutorial. If you've tweaked parameters for that, you can also update them in MainActivity to get the same results as in your server testing.

[Download this project from GitHub](#)

Related Post

[Android TensorFlow Machine Learning](#)

[Google Cloud Speech API in Android APP](#)

Share this post:

[on Twitter](#)

[on Facebook](#)

[on Google+](#)

Previous Article

< [EmojiCompat Support Library](#)

Next Article

[How to Create Instant app from Existing App](#) >

Leave a Reply

Your email address will not be published. Required fields are marked *

Comment

Name *

Email *

Website

Post Comment

Recent Posts

- [How Location API Works in android](#)
- [Firebase Crashlytics](#)
- [How to use DateTime datatype in SQLite Using Room](#)
- [Cloud Firestore Database for Web Application](#)
- [ConstraintLayout 1.1.0: Circular Positioning](#)
- [Firestore Security With Firebase Authentication](#)
- [Firestore Document Database Data Model](#)
- [Firebase Firestore Database for Android Application](#)
- [Rest API Pagination with Paging Library.](#)
- [Architecture Components:Paging Library](#)
- [Autosizing TextViews Using Support Library 26.0](#)
- [How to Create Instant app from Existing App](#)
- [Speech Recognition Using TensorFlow](#)
- [EmojiCompat Support Library](#)
- [Fast Scrolling in RecyclerView Using Support Library 26](#)

Categories

- [architecture \(7\)](#)
- [Connectivity \(2\)](#)
- [Firebase \(8\)](#)
- [kotlin \(1\)](#)
- [Layout \(1\)](#)
- [Library \(1\)](#)
- [Location \(2\)](#)
- [Machine Learning APIs \(3\)](#)
- [Performance \(2\)](#)
- [TensorFlow \(6\)](#)
- [Uncategorized \(13\)](#)

Archives

- [December 2017](#)
- [November 2017](#)
- [October 2017](#)
- [September 2017](#)
- [August 2017](#)
- [July 2017](#)
- [June 2017](#)
- [May 2017](#)
- [April 2017](#)

Proudly powered by [androidkt](#)