




Images haven't loaded yet. Please exit printing, wait for images to load, and try to print again.

and waiting for my own creativity to kick in!  
Nov 26, 2016 · 4 min read



Official TensorFlow background

- **TensorFlow: How to freeze a model and serve it with a python API**

We are going to explore two parts of using a ML model in production:

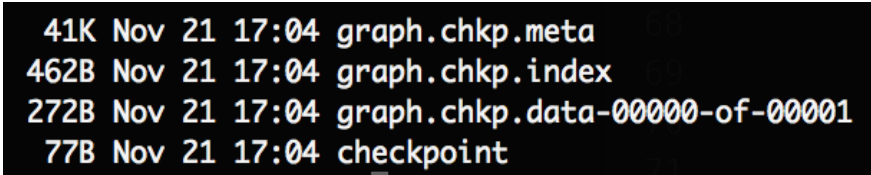
- How to export a model and have a simple self-sufficient file for it
- How to build a simple python server (using flask) to serve it with TF

*Note: if you want to see the kind of graph I save/load/freeze, you can here*

### **How to freeze (export) a saved model**

*If you wonder how to save a model with TensorFlow, please have a look at my previous article before going on.*

let's start from a folder containing a model, it probably looks something like this:



```
41K Nov 21 17:04 graph.chkp.meta
462B Nov 21 17:04 graph.chkp.index
272B Nov 21 17:04 graph.chkp.data-00000-of-00001
77B Nov 21 17:04 checkpoint
```

Screenshot of the result folder before freezing our model

The important files here are the “**.chkp**” ones. If you remember well, for each pair at different timesteps, one is holding the weights (“**.data**”) and the other one (“**.meta**”) is holding the graph and all its metadata (so you can retrain it etc...)

But when we want to serve a model in production, we don’t need any special metadata to clutter our files, we just want our model and its weights nicely packaged in one file. This facilitate storage, versioning and updates of your different models.

Luckily in TF, we can easily build our own function to do it. Let’s explore the different steps we have to perform:

Retrieve our saved graph: we need to load the previously saved meta graph in the default graph and retrieve its graph\_def (the ProtoBuf definition of our graph)

Restore the weights: we start a Session and restore the weights of our graph inside that Session

Remove all metadata useless for inference: Here, TF helps us with a nice helper function which grab just what is needed in your graph to perform inference and returns what we will call our new “**frozen graph\_def**”

Save it to the disk, Finally we will serialize our frozen graph\_def ProtoBuf and dump it to the disk

Note that the two first steps are the same as when we load any graph in TF, the only tricky part is actually the graph “**freezing**” and TF has a built-in function to do it!

I provide a slightly different version which is simpler and that I found handy. The original freeze\_graph function provided by TF is installed in your bin dir and can be called directly if you used PIP to install TF. If

not you can call it directly from its folder (see the commented import in the gist).

So let's see:

```
1  import os, argparse
2
3  import tensorflow as tf
4
5  # The original freeze_graph function
6  # from tensorflow.python.tools.freeze_graph import freeze_graph
7
8  dir = os.path.dirname(os.path.realpath(__file__))
9
10 def freeze_graph(model_dir, output_node_names):
11     """Extract the sub graph defined by the output node names
12     all its variables into constant
13
14     Args:
15         model_dir: the root folder containing the checkpoints
16         output_node_names: a string, containing all the output node names,
17                             comma separated
18     """
19     if not tf.gfile.Exists(model_dir):
20         raise AssertionError(
21             "Export directory doesn't exists. Please supply a valid"
22             "directory: %s" % model_dir)
23
24     if not output_node_names:
25         print("You need to supply the name of a node to freeze")
26         return -1
27
28     # We retrieve our checkpoint fullpath
29     checkpoint = tf.train.get_checkpoint_state(model_dir)
30     input_checkpoint = checkpoint.model_checkpoint_path
31
32     # We precise the file fullname of our freezed graph
33     absolute_model_dir = "/"
34     output_graph = absolute_model_dir + "/frozen_model.pb"
35
36     # We clear devices to allow TensorFlow to control more devices
37     clear_devices = True
38
39     # We start a session using a temporary fresh Graph
40     with tf.Session(graph=tf.Graph()) as sess:
41         # We import the meta graph in the current default graph
42         saver = tf.train.import_meta_graph(input_checkpoint, sess.graph)
```

Now we can see a new file in our folder: **“frozen\_model.pb”**.



Screenshot of the result folder after freezing our model

As expected, its size is bigger than the weights file size and lower than the sum of the two checkpoints files sizes.

Note: In this very simple case, the weights size is very small, but it is usually multiple Mbs.

## How to use the frozen model


Naturally, after knowing how to freeze a model, one might wonder how to use it.

The little trick to have in mind is to understand that what we dumped to the disk was a graph\_def ProtoBuf. So to import it back in a python script we need to:


Import a graph\_def ProtoBuf first

Load this graph\_def into a actual Graph

We can build a convenient function to do so:



Now that we built our function to load our frozen model, let's create a simple script to finally make use of it:




Note: when loading the frozen model, all operations got prefixed by “prefix”. This is due to the parameter “name” in the “import\_graph\_def” function, by default it prefix everything by “import”.

This can be useful to avoid name collisions if you want to import your graph\_def in an existing Graph.

## How to build a (very) simple API

For this part, I will let the code speaks for itself, after all this is a TF series about TF and not so much about how to build a server in python. Yet it felt kind of unfinished without it, so here you go, the final workflow:

*Note: We are using flask in this example*



. . .

## TensorFlow best practice series

This article is part of a more complete series of articles about TensorFlow. I've not yet defined all the different subjects of this series, so if you want to see any area of TensorFlow explored, add a comment! So far I wanted to explore those subjects (this list is subject to change and is in no particular order):

- A primer
- How to handle shapes in TensorFlow
- TensorFlow saving/restoring and mixing multiple models
- How to freeze a model and serve it with python (this one!)
- TensorFlow: A proposal of good practices for files, folders and models architecture
- TensorFlow howto: a universal approximator inside a neural net
- How to optimise your input pipeline with queues and multi-threading
- Mutating variables and control flow
- How to handle input data with TensorFlow.
- How to control the gradients to create custom back-prop or fine-tune my models.
- How to monitor my training and inspect my models to gain insight about them.

**Note:** *TF is evolving fast right now, those articles are currently written for the 1.0.0 version.*

. . .

## References

- [https://github.com/tensorflow/tensorflow/blob/master/tensorflow/python/tools/freeze\\_graph.py](https://github.com/tensorflow/tensorflow/blob/master/tensorflow/python/tools/freeze_graph.py)





